

POLITECNICO DI TORINO

EURECOM

Master's Degree in Data Science and Engineering



Master's Degree Thesis

Personal Data Detection in Free Text

Supervisors:

Dr. Francesco Di Cerbo

Prof. Paolo Papotti

Prof. Giuseppe Rizzo

Candidate:

Gabriele Gioetto

Academic Year 2022/2023
Torino, Sophia Antipolis

Abstract

An HR (Human Resources) department in a large organization receives inquiries/requests from employees on multiple topics, which are quite different from one another. As an example, an employee can send requests dealing with health conditions, compensation/taxation, events of life (marriage, death of a relative...).

These data can be used for many different queries that can be useful for analysis purposes (Example: ‘How many people have had COVID during 2021’). However, HR tickets typically contain personal data, that cannot be processed without the consent of the data subject according to the European privacy regulation (GDPR).

To be able to process documents with personal data, we can identify the pieces of information that qualify as personal data in a communication and subsequently anonymize such information using the appropriate techniques. A significant part of this problem is represented by the complex nature of personal data according to GDPR: personal data are defined as ‘*any piece of information that can be connected to an identified or identifiable natural person*’. They comprise obvious identifiers like social security numbers, email addresses, but also elements like ‘the Italian intern working for SAP in South of France’. To the best of our knowledge, it does not exist a public dataset of HR tickets that can be used to train machine learning models, the main reason being the difficult nature of these types of data. Synthetic data, which are artificial data that are generated from original data using a model that is trained to reproduce the characteristics and structure of the original data, follow a data protection by design approach. To address the need for a large dataset of HR tickets, we created a taxonomy of tickets, we found real data that can be used as support to create synthetic tickets and developed *Ticket Generator*: an application that can produce as many tickets as needed

belonging to different categories, we released a dataset of previously created tickets and we showcase some possible use cases of the dataset.

Abstract French

Le service RH (Ressources Humaines) d’une grande entreprise reçoit des demandes de renseignements de la part des employés sur de multiples sujets, très différents les uns des autres. Par exemple, un employé peut envoyer des demandes concernant l’état de santé, la rémunération/imposition, les événements de la vie (mariage, décès d’un proche...).

Ces données peuvent être utilisées pour de nombreuses requêtes différentes qui peuvent être utiles à des fins d’analyse (Exemple : ‘Combien de personnes ont eu un COVID au cours de l’année 2021’). Cependant, les tickets RH contiennent généralement des données personnelles, qui ne peuvent être traitées sans le consentement de la personne concernée, conformément au Règlement Général sur la Protection des Données (RGPD).

Pour pouvoir traiter les documents contenant des données à caractère personnel, nous pouvons identifier les éléments d’information qui sont qualifiés de données à caractère personnel dans une communication et, par la suite, rendre ces informations anonymes en utilisant les techniques appropriées. Une partie importante de ce problème est représentée par la nature complexe des données personnelles selon le RGPD: les données personnelles sont définies comme ‘*toute information se rapportant à une personne physique identifiée ou identifiable*’. Elles comprennent des identifiants évidents comme les numéros de sécurité sociale, les adresses e-mail, mais aussi des éléments comme “le stagiaire italien travaillant pour SAP dans le sud de la France“. À notre connaissance, il n’existe pas de base de données publique des tickets RH qui peuvent être utilisée pour entraîner des modèles d’apprentissage automatique, principalement en raison de la nature difficile de ces types des données. Les données synthétiques, qui sont des données artificielles générées à partir de données originales et d’un modèle entraîné à reproduire

les caractéristiques et la structure des données originales, suivent une approche de protection des données par conception. Pour répondre au besoin d'un grand ensemble de données de tickets RH, nous avons créé une taxonomie de tickets, nous avons trouvé des données réelles qui peuvent être utilisées comme support pour créer des tickets synthétiques et nous avons développé *Ticket Generator*: une application qui peut produire autant de tickets que nécessaire appartenant à différentes catégories, nous avons publié un ensemble de tickets précédemment créés et nous présentons quelques cas d'utilisation possibles de l'ensemble de données.

Acknowledgements

About SAP

Founded in 1972, SAP has grown to become the world's leading provider of business software solutions. SAP is market leader in enterprise application software. The company is also the fastest-growing major database company. Globally, more than 77% of all business transactions worldwide touch an SAP software system.

With more than 347.000 customers in more than 180 countries, SAP includes subsidiaries in all major countries. SAP is the world's largest inter-enterprise software company and the world's third-largest independent software supplier, overall. SAP solutions help enterprises of all sizes around the world to improve customer relationships, enhance partner collaboration and create efficiencies across their supply chains and business operations. SAP employs more than 98.600 people.

Security Research at SAP Labs France, Sophia Antipolis

Based at SAP Labs France Mougins, Security Research Sophia-Antipolis[1] addresses the upcoming security needs, focusing on increased automation of the security life cycle and on providing innovative solutions for the security challenges in networked businesses, including cloud, services and mobile. This internship is based in the SAP Labs France Research Lab, in Sophia-Antipolis. The work will be performed in the context of the Research Program "Security & Trust".

Table of Contents

List of Figures	XII
Acronyms	XV
1 Introduction	1
1.1 GDPR	1
1.2 Synthetic data	2
1.3 Ticket Generation	3
1.4 Structure of the thesis	3
2 Related Works	6
3 Method	8
3.1 Taxonomy	8
3.2 Datasets	10
3.2.1 Datasets preprocessing	11
3.2.2 Bayesian Networks	13

3.3	Ticket Generation	15
3.4	State of The Art auto-regressive language model	20
3.5	GPT-J	21
3.5.1	Parameters	22
3.5.2	Architecture	24
3.6	Architecture analysis	30
3.7	Survey	35
3.8	Results	36
3.9	Additional Experiments	42
3.9.1	Topical Generation	43
3.9.2	Fine-tune GPT	45
4	Use Cases	46
4.1	Classification	47
4.1.1	FastText Classifier	47
4.1.2	BERT classifier	48
4.1.3	Results	49
4.2	Anonymization	52
4.2.1	QAGS	55
4.2.2	SUMMAQA	56
4.2.3	Cosine Similarity	57
4.3	Named Entity Recognition	60

4.3.1	Weak labeling	60
4.3.2	Classical NER	62
4.3.3	Our approach to NER	63
5	Ethical Considerations	68
6	Conclusion	69
	Bibliography	71

List of Figures

1.1	Schema of the generation of the dataset	4
1.2	Schema of the possible use cases of the dataset	5
3.1	Initial Taxonomy	9
3.2	Final taxonomy, reduced due to unavailability of data	10
3.3	Bayesian network dependencies	12
3.4	Schema of Ticket Generation	18
3.5	GPT-3 and GPT-J architectures compared	27
3.6	Byte-Pair Encoding Tokenization Example	28
3.7	Implementation of RoPE, Image taken from original paper[20]	30
3.8	Input Saliency: First token	32
3.9	Input Saliency: Subject	32
3.10	Input Saliency: HR	33
3.11	Non-negative Matrix Factorization on Activation Matrix . .	33
3.12	Neuron Activation Analysis	34
3.13	Example of survey	36

3.14	Hierarchy of all topics generated by BERTopic	42
3.15	Main 8 topics generated by BERTopic	43
4.1	Schema of BERT classification	49
4.2	Confusion matrix of Ticket classification with BERT	51
4.3	Schema of Ticket Anonymization	54
4.4	Schema of QAGS	56
4.5	Schema of SUMMAQA	58
4.6	Schema of Cosine Similarity for	59
4.7	Ticket NER preprocessing part 1	65
4.8	Ticket NER preprocessing part 2	65
4.9	Ticket NER training	66
4.10	Ticket NER confusion matrix	67

Acronyms

AI

Artificial Intelligence

BERT

Bidirectional Encoder Representations from Transformers

BLEU

BiLingual Evaluation Understudy

CPU

Central Processing Unit

EU

European Union

GDPR

General Data Protection Regulation

GPT

Generative Pre-trained Transformer

GPU

Graphics Processing Unit

HR

Human Resources

ICD

International Classification of Diseases

LLM

Large Language Model

LM

Language Model

LSA

Latent Semantic Analysis

MLM

Masked Language Modeling

NER

Named Entity Recognition

NMF

Non-negative Matrix Factorization

NSP

Next Sentence Prediction

PII

Personally Identifiable Information

POS

Part Of Speech

RAM

Random Access Memory

ROUGE

Recall-Oriented Understudy for Gisting Evaluation

SVD

Singular Value Decomposition

T5

Text-To-Text Transfer Transformer

TF-IDF

Term Frequency–Inverse Document Frequency

TPU

Tensor Processing Unit

TTR

Type Token Ratio

Chapter 1

Introduction

With the advent of deep learning, a lot of applications have been needing more and more data. However, more often than not, these data contain a large amount of sensitive and personal information that restricts its use according to the legal framework in place in many countries.

Data is growing rapidly within companies, containing a great deal of sensitive information that can't be processed or shared without due consideration of legal ramifications. If such data could reveal the identity of someone, their personal rights would be threatened.

This is especially relevant in the context of HR tickets, which can contain not only personal information but also special categories of personal data, which require extra protection according to GDPR.

1.1 GDPR

The General Data Protection Regulation[2] is a privacy regulation that regulates the processing of personal data *"wholly or partly by automated means and to the processing other than by automated means"*. The regulation applies to all citizens of the EU and to all data subjects in the EU, whether the processing is carried out inside or outside the EU.

Personal data are defined as *"any information relating to an identified or*

identifiable natural person". The regulation states that *"an identifiable natural person is one who can be identified, directly or indirectly, in particular by reference to an identifier such as a name, an identification number, location data, an online identifier or to one or more factors specific to the physical, physiological, genetic, mental, economic, cultural or social identity of that natural person"*. The natural person can be identified both by direct identifiers and quasi-identifiers. The direct identifiers are information that directly identify the person, such as the telephone number, social security number. . . , while the quasi-identifiers are information that alone cannot identify a person, but if they are combined with other quasi-identifiers can affect a person's privacy. For example, the job title could be not enough to identify a person, but combined with his/her company and his/her nationality it could be. Moreover, the GDPR treats some categories of personal data more carefully. These categories are called 'special categories' and include racial or ethnic origin, political opinions, religious or philosophical beliefs, trade union membership, genetic data, biometric data, data concerning health or data concerning a natural person's sex life or sexual orientation. The special categories of personal data cannot be generally processed, with some exceptions including "the data subject has given explicit consent to the processing of those personal data for one or more specified purposes".

1.2 Synthetic data

Synthetic data is artificial data that is generated from real data and has the same statistical distribution as the original data . This means that synthetic data and original data should deliver very similar results when undergoing the same statistical analysis.

Synthetic data has many benefits over real data: if you create a model that generates synthetic data you can generate how many data you need, you can infer certain properties to your data (for example it can be useful for bias and fairness research) and above all, synthetic data can respect the right to personal data protection. It is important to point out that it is not always guaranteed that synthetic data is privacy-preserving: it has been shown that synthetic data can leak personal information [3].

1.3 Ticket Generation

We created a tool that can be used to create an unlimited amount of synthetic HR tickets and we published a dataset of 16000 tickets. Each ticket, other than the ticket's text, is composed of a category, sub-category and the entities of the ticket. The tickets are not created from scratch, but starting from real-world's datasets and some prompts that help the model generate the text.

We did a survey internal to the company to gather some real data, and we changed the ticket generation parameters in order to respect the real data with respect to different text metrics.

Finally, we showed different possible use cases for the dataset:

- Ticket Anonymization
- Ticket Classification
- Named Entity Recognition on tickets' entities

1.4 Structure of the thesis

The thesis is divided in two main parts: the dataset generation (chapter 3 and Figure 1.1) and the possible use cases of the dataset (chapter 4 and Figure 1.2).

First, we built a taxonomy of the tickets (section 3.1), then we listed all the datasets which we used as a starting point to get some initial data and we explained how they were processed (section 3.2). After that, we described the process of the generation of the tickets (section 3.3), where we showed the templates (section 3.3), the generative model (section 3.5) and the parameters we can tune during the generation process (subsection 3.5.1). Then, we did some analysis on the generation outcomes, both on the model (section 3.6) and on the tickets' texts (section 3.8). In section 3.7 we illustrate the survey we did to gather the test dataset.

Lastly, in section 3.9 we introduced some experiments we carried out but

that were not included in the final version of the application.

In chapter 4 we listed the three different use cases for the dataset and their results: the classification of the tickets' categories (section 4.1), the anonymization of the tickets with a different approach than the standard one (section 4.2) and the named entity recognition on the entities of the tickets (section 4.3), where we used both a classical approach (subsection 4.3.2) and an innovative one (subsection 4.3.3).

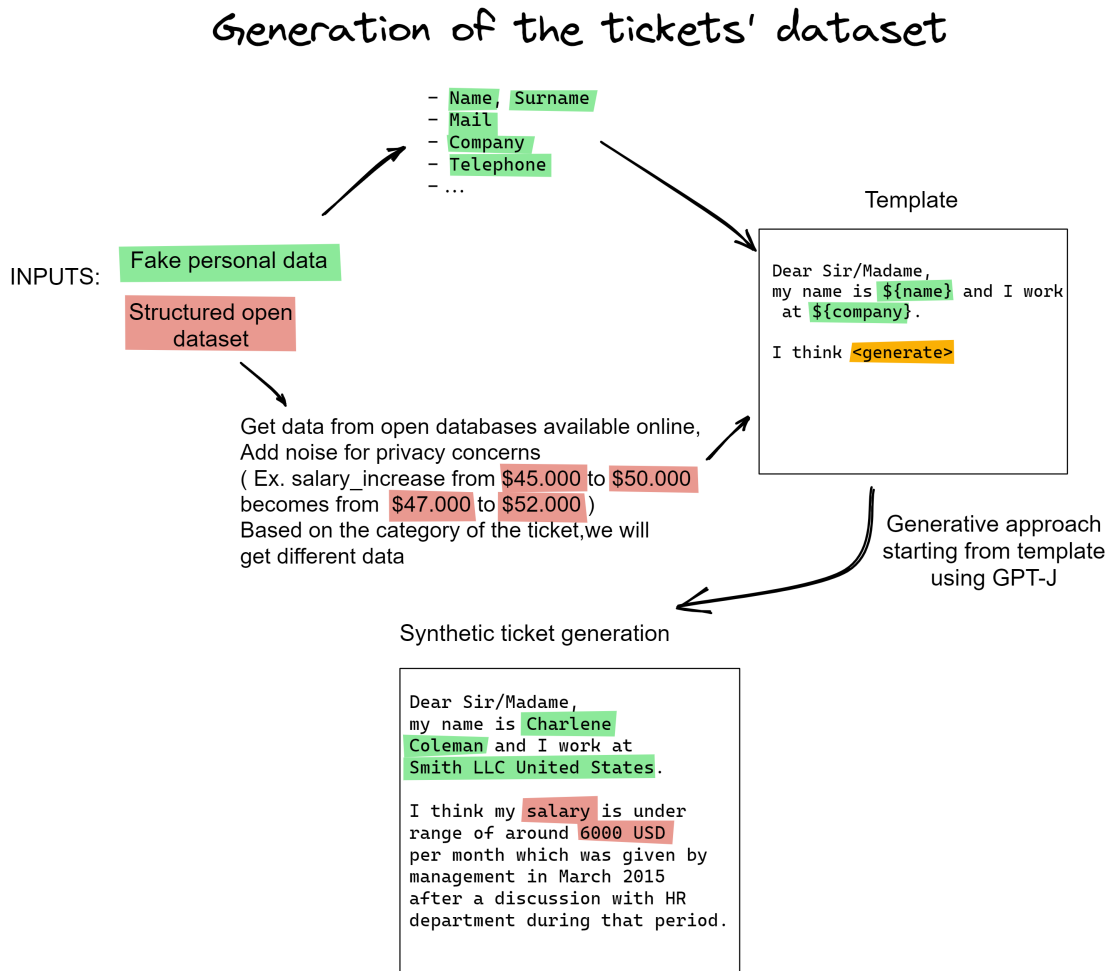
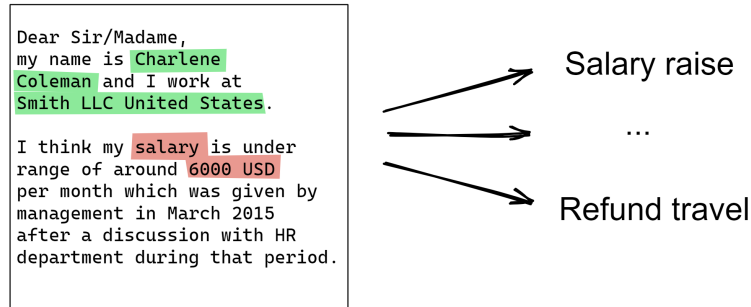


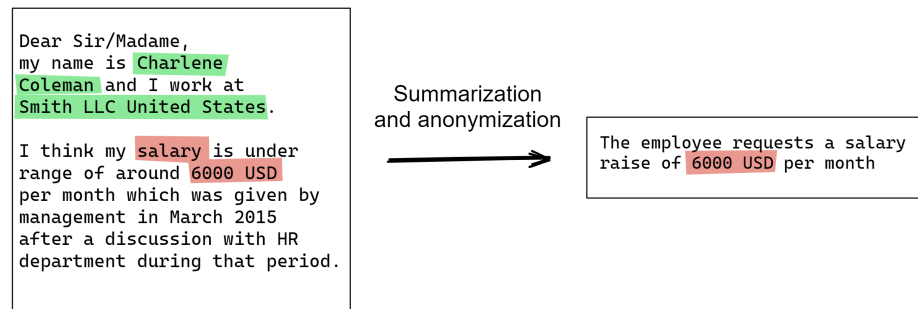
Figure 1.1: Schema of the generation of the dataset

USE CASES

Classification



Anonymization



Named Entity Recognition

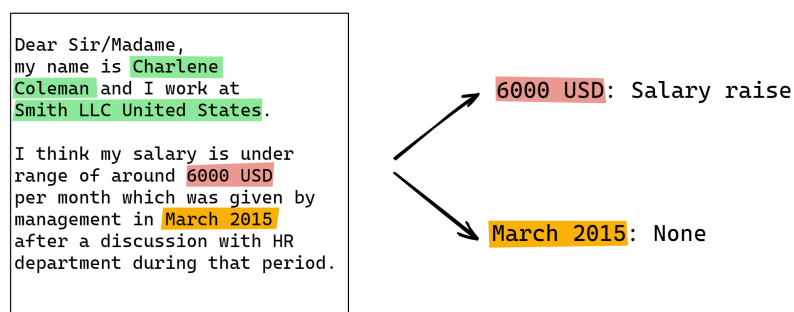


Figure 1.2: Schema of the possible use cases of the dataset

Chapter 2

Related Works

Due to the sensitive nature of the data and GDPR laws, a dataset of HR tickets written by company employees has never been published as far as we know. This is likely to be the case because the public release of these data could be damaging for both the company and the employees, such as an employee criticizing the working environment or revealing confidential information about the company in a ticket.

However, some datasets can resemble what we are trying to build, with similar language style and intents:

- *HR Analytics*: data about 1470 employees in a company: the features of the dataset include job role, gender, education field, hourly rate. . .
- *Consumer Complaint*: Consumer Financial Protection Bureau’s online database of customer complaints about different financial products. The overall dataset contains over 600,000 complaints and each record has the complaint’s text description, the product that is the cause of the complaint, the company which issues the product, and the category of the issue. In the dataset the personal information are masked.
- *Enron Mail*: The *Enron Mail* dataset contains about 0.5M emails of 150 senior management from the Enron corporation. This data was originally made public, and posted to the web, by the Federal Energy Regulatory Commission during its investigation for fraud. The corpus

is one of the only few publicly available mass collections of real emails easily available for study.

- *TAB*: the Text Anonymization Benchmark corpus[4] is a privacy-oriented annotated text resource. The corpus comprises 1,268 English-language court cases from the European Court of Human Rights (ECHR) enriched with comprehensive annotations about the personal information appearing in each document, including their semantic category, identifier type, confidential attributes, and co-reference relations.

Chapter 3

Method

3.1 Taxonomy

When building a dataset is important to decide the features that will be used. Features should be chosen carefully to ensure that the dataset is relevant and useful. In most HR ticketing systems, the tickets belong to a category, which helps the HR department to classify the tickets and send them to the right person to answer them.

Usually, HR tickets can belong to various categories, which can range from a request for a shift change to a complaint about a colleague.

We built a taxonomy of tickets, which is structured into categories and subcategories. The categories and subcategories let us define a finite set of possible topics of the tickets and give a precise structure to the dataset.

Each subcategory has its own variables that are used as inputs for the tickets' generation. The variables are sampled from real-world datasets.

For example, to create a request for sick leave, we pass as inputs to our model the reason and the number of days of sick leave requested, which will be acquired from an external dataset.

We have a selection of templates and prompts to kick-start the generation process. Every category has its own distinct templates and prompts. The taxonomy (Shown in Figure 3.1) has been built inspired by industrial approaches on the categorization of most common HR requests, however, the final taxonomy (Shown in Figure 3.2) presented here is a subset of the

original one due to the unavailability of public data on certain topics (Ex. *Work benefits*)

The final complete taxonomy and the complete list of all variables used for each category/sub-category can be seen in the Table 3.1

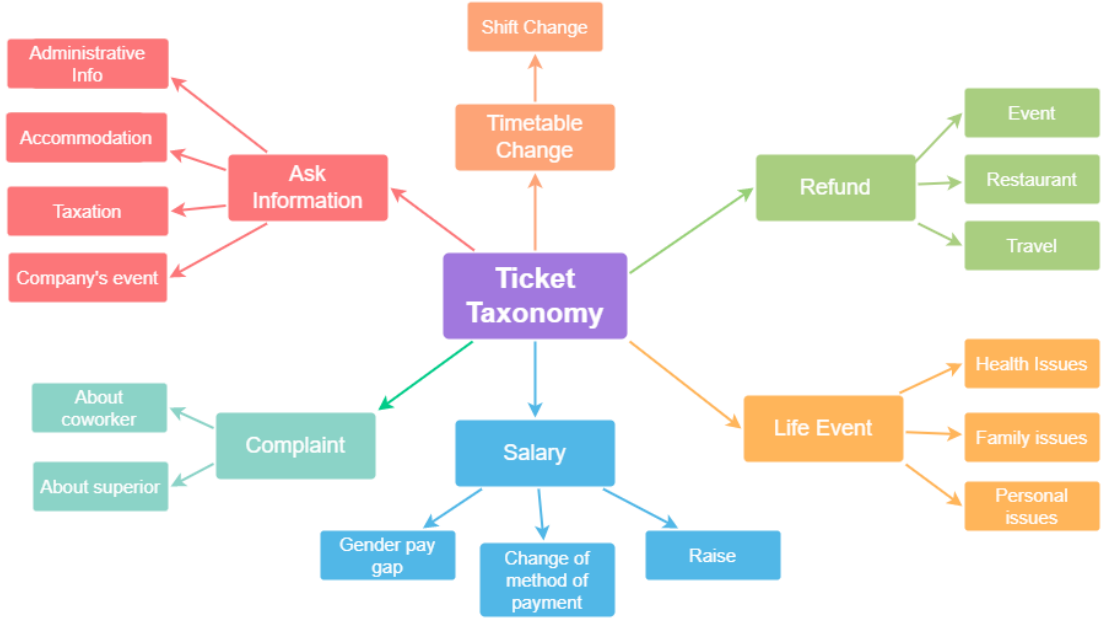


Figure 3.1: Initial Taxonomy

Category	Sub-category	variables
Ask Information	Accommodation	location, duration
Complaint	About coworker	complaint, reason
	About superior	complaint, reason
Timetable change	Shift change	reason_of_change, old_date, new_date
Salary	Salary raise	old_salary, new_salary, increase, work_title
	Gender pay gap	wage_gap
Life Event	Health issues	disease, number_of_days_of_sick_leave
	Personal issues	issue, number_of_days
Refund	Travel	from, to, date_travel

Table 3.1: Table of all defined categories and sub-categories with their respective variables




Figure 3.2: Final taxonomy, reduced due to unavailability of data

3.2 Datasets

In order to generate tickets, we decided to use real data as a starting point to make them as much realistic as possible. Another reason to use real data is that it makes the dataset useful for use cases such as anonymization. The datasets that we used are all public and available online. In some cases where no datasets were available, we created them manually from personal experience.

The datasets are:

- *Absenteeism at work Data Set*: contains records of work absences, with the reason for the absence (almost always a disease) and the number of hours of absence. It is used to create requests for days off for health reasons. The data were gathered from July 2007 to July 2010 at a courier company in Brazil.
- *National Occupational Employment and Wage Estimates United States*: estimates of wages in the US calculated with data collected from employers in all industry sectors in metropolitan and nonmetropolitan areas in

every state and the District of Columbia. It is used to create requests for salary raises.

- *List of events of life*: list of major events in life. It is used to create requests for time off due to personal reasons.
- *Gender pay gap in the UK*: dataset of employers with 250 or more employees, comparing men’s and women’s average pay across the organizations in 2021. It is used to create the requests for explanations for the wage gap between genders.
- *OpenFlights database*: datasets of airports and flights all over the world. It is used for requests for refunds of travel. The data are updated up to 2017.
- *Geonames all cities with a population over 1000*: datasets of all cities of the world with a population over 1000 people. It is used for requests for information about accommodation.

3.2.1 Datasets preprocessing

Absenteeism at work Data Set

The *Absenteeism at work Data Set* is distinct because it is the only dataset that contains individual records of employees, not just a collective average. These data include the personal details of the employee, the cause of their absence, and the duration of the absence in hours. To ensure the privacy of the subjects, we utilized a Bayesian Network.

A Bayesian network is a probabilistic graphical model that measures the conditional dependence structure of a set of random variables based on the Bayes theorem. The features that we have used to build the Bayesian network are the *reason for absence*, the *month of absence* and the *absenteeism time*. We added the dependency between *month of absence* and *reason for absence*, since some illnesses are more common during some months of the year, and between *absenteeism time* and *reason for absence*, since based on the type of illness you can request more or less time off (Figure 3.3).

Using the *Absenteeism at work Data Set* we learn conditional probability distributions from data, to which we add a Laplace noise for differential




Figure 3.3: Bayesian network dependencies

privacy. Then we can sample new data that follow the original distributions, but that are not equal to the original ones.

The absence reasons in the dataset are given as ICD (International Classification of Diseases) codes, to make them more human-readable, we picked for each ICD code the corresponding more frequent diseases.

National Occupational Employment and Wage Estimates United States

In the *National Occupational Employment and Wage Estimates United States* dataset, we sample employees based on the number of people employed in a certain field. Therefore for example since ‘Retail Sales Workers’ consists of 5.4% of the total occupations, then the sampled employee will have the 5.4% of possibility to have as occupation ‘Retail Sales Worker’.

The current salary of the employee is calculated by adding Gaussian noise to the average salary of the employee’s occupation, and then the salary raise requested is picked randomly between 5%-10%.

Gender pay gap in the UK

The ranges of the wage gap, used in the tickets regarding the explanation for the gender wage gap in the company, are sampled randomly from the dataset *Gender pay gap in the UK*, adding a Gaussian noise for privacy reasons.

Geonames all cities with a population over 1000

To sample the cities for the requests of accommodation, we randomly sample from all the cities with more than 100,000 inhabitants from the country of residence of the synthetic employee. To calculate the duration of the accommodation we pick a random number of months between 1 and 12.

OpenFlights database

To get data for the category type "refund travel", we sample randomly one flight from all the flights leaving from the country of the synthetic employee. The data are taken from the *OpenFlights database*.

Complaints - life events

The complaints about coworkers and superiors and the life events that can affect the work life of a person were written by myself, using primarily personal experience and some help from internet blogs.

3.2.2 Bayesian Networks

A Bayesian network is a machine learning method that combines a probabilistic graphical model with Bayesian inference to infer the likelihood of certain events or outcomes. It is used to find relationships between variables and to identify which variables are most influential in predicting a certain outcome or event.

The bayesian network is based on the Bayesian theorem:

$$p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$

Formally, the Bayesian network is a directed acyclic graph $G = (V, E)$ with

- A feature for each node i belonging to V
- A conditional probability distribution for each edge, so the edge from feature i to j represents $p(x_j | x_i)$

The base version of a Bayesian network works with discrete variables, however, some implementations consider also continuous variables [5].

To have privacy-preserving data, which is data that does not include any personal information that could identify an individual, we inject some noise into the Bayesian Network, exploiting a PrivBayes[6].

PrivBayes is a method for releasing high-dimensional data while ensuring differential privacy. It begins by constructing a Bayesian network that compresses the correlations between the attributes in the dataset D , allowing us to approximate the distribution of data with a set P of low-dimensional marginals. Noise is then injected into each marginal in P to uphold the differentially private guarantee, and the Bayesian network and noisy marginals are used to create an approximation of the data distribution in D . Tuples are sampled from this approximate distribution to construct a synthetic dataset. In practice implementing PrivBayes for our case is relatively easy: we build a Bayesian network starting from the *Absenteeism at work Data Set* calculating the likelihood distribution $p(x_i | x_j) \forall x_i, x_j \in D$, where D is the set of features. As a prior, we used a Dirichlet distribution, mainly because it is the conjugate prior of the categorical distribution.

We then added pseudo counts to the observed counts in the data used to calculate $p(x_j | x_i)$. This technique is used to diminish the overfitting of data. The values we used for pseudocount is $\gamma = 1$.

Since we learn the conditional probability distribution from our data, the structure of the network or the conditional probabilities may therefore leak some information on an individual in the dataset. In order to provide strong privacy guarantees and minimize the re-identification risk, we leverage the

notion of differential privacy: we perturb the data adding a noise sampled from a Laplace distribution

$$z \sim \text{Laplace}\left(0, \frac{2 \cdot n_{\text{features}}}{\gamma \cdot \epsilon}\right)$$

where ϵ is the privacy budget for differential privacy, which controls the anonymization level.

Differential privacy is a rigorous mathematical definition of privacy. An algorithm is said to be differentially private if by looking at the output of an algorithm A performed on a dataset X , one cannot tell whether any individual's data was included in the dataset X or not. In other words, a differentially private algorithm guarantees that its behavior hardly changes when a single individual joins or leaves the dataset. The mathematical definition of differential privacy is:

$$\Pr[A(X) \in Z] \leq e^\epsilon \cdot \Pr[A(X') \in Z]$$

where A is an algorithm and X' is a neighbour dataset of X . A dataset is a neighbor of another dataset if they differ by only one record.

Once the private Bayesian network is built, we can sample new values for all the nodes in the graph. These generated values will have the same distribution and preserve the consistency and statistical properties of the original dataset up to the noise addition which acts as a de-identification barrier.

3.3 Ticket Generation

For each HR ticket, we create a synthetic employee. Depending on the category of the ticket we want to generate, the employee can have different features. For all tickets' categories, the employees have some common features: *name*, *first name*, *last name*, *nationality*, *country*, *email*, *company*, *company's email* and *ticket date*.

All these information are created exploiting the Python library *Faker*[7]. The nationality and the company's country are selected from the extendible list { USA, Germany, Italy, Spain, France }. All other information are created accordingly to the country picked. So for example, if the country of birth of

the employee is Italy, then the generated name will be Italian.

Then, once the employees are generated, the information specific to the ticket category, the ones created starting from the open datasets as mentioned before, are concatenated to the general information of the employees.

For each ticket category, there are distinct templates. In each template there is an initial part that contains the general information of the employee, such as name, surname, company. . . , then some prompts correlated to the category of the ticket and then the textual prompt.

Here are a couple of examples of templates:

Request for time off due to health reasons:

From: \${email}
To: \${company email}
First name: \${first name}
Last name: \${last name}
Company: \${company}
Date: \${ticket date}
Ticket category: \${category}
Ticket sub-category: \${sub category}
Date start absence: \${date start absence}
Reason absence: \${reason}
Subject: Request for sick leave for \${number of days}

Dear Sir/Madame, my name is \${name} and I work at \${company}. I am requesting *<generate>*. I hope *<generate>*.

Request for refund of travel:

From: \${email}
To: \${company email}
First name: \${first name}
Last name: \${last name}
Company: \${company}
Date: \${ticket date}
Ticket category: \${category}
Ticket sub-category: \${sub category}

Date Travel: \${date travel}
From: \${airport from}, \${from}
Destination: \${airport to}, \${to}
Subject: Request of refund for plane ticket from \${from} to \${to} at
\${date travel}
Hello, my name is \${name}. I am writing this mail to ask for a refund
for the travel <generate>

The variables are replaced with the features of the employee, whereas the <generate> are replaced with text generated by a generative model. The parts generated by the generative model are created recursively. This means that the first <generate> is replaced with text generated automatically using as prompt everything that precedes it. Then the second <generate> will have as a prompt the entire ticket, including the text generated previously by the model. The model is forced to generate some text, if no text is generated in an iteration, the process is repeated until the model gives a non-empty output.

A general schema of the generation is shown in Figure 3.4

Templates

Generative Pre-trained Transformer (GPT) [8] models take in a prompt, or context, and generate text from it. The prompts typically take the form of a few sentences or a paragraph, and the models generate sentences that fit the context of the prompt. By manipulating the prompt, users can generate text of various tones, topics, and styles.

The GPT models are also capable of completing tasks such as question answering, machine translation, and summarization in an unsupervised manner[9]. By providing a prompt with the task and context, the models can generate accurate results that address the specific context and task. For example, summarization tasks require a prompt to provide the necessary data so that the model can deliver a correct summary of the text.

Changing the prompt of GPT can change the tone, topics, and style of the generated text. Depending on the prompt, GPT models can generate text ranging from creative stories to technical summaries. The tone and style of




Figure 3.4: Schema of Ticket Generation

the text can range from humorous to academic, depending on the prompt. As the prompt changes, the model will also adjust to reflect the context of the prompt.

This is why we took inspiration from the e-mail format, as HR tickets are created in a working environment where formal language is used and often they resemble emails in the tone and the topics.

The emails of the *Enron dataset*, one of the few datasets of emails public available, have usually a well-structured prompt, here's an example:

Message-ID: <16593073.1075858228177.JavaMail.evans@thyme>
Date: Wed, 12 Jan 2000 00:29:00 -0800 (PST)
From: carrie.hollomon@enron.com
To: phillip.love@enron.com

Subject: Workhours
Mime-Version: 1.0
Content-Type: text/plain; charset=us-ascii
Content-Transfer-Encoding: 7bit

Hello ...

To mimic the format of the emails of the *Enron dataset*, we kept the From, To, Date and Subject rows. Instead, we removed the Mime-Version, the Content-type and Content-Transfer-Encoding, because after conducting some experiments it was evident that they did not help to achieve better results, on the contrary in some cases they were worse.

In addition to the standard information, we added some rows with additional information specific to the category, rather than including them only in the subject or in the initial prompt.

In order to generate tickets that were dissimilar and covered a wide range of topics/tokens, we preferred giving GPT small text prompts and letting the model generate most of the text getting the information from the email-like prompt. This approach allowed us to improve the diversity of the dataset and, consequently, its usefulness.

Small text prompt example of a request for shift change:

Dear Sir/Madame, my name is \${name}. I wanted to <generate>

Long text prompt example of a request for shift change:

Dear Sir/Madame, my name is \${name} and I work at \${company}. I wanted to ask to change the shift from \${old_date} to \${new_date} in order to be able to <generate >

3.4 State of The Art auto-regressive language model

Autoregressive transformer-based Large Language Models are sequence-to-sequence deep learning models that are pre-trained on a large corpus of data. They are designed to generate new tokens conditioning the model on some input text. The output probability distribution for the next token that the model generates is a probability distribution over all possible tokens in the corpus.

Autoregressive transformer-based Large Language Models have shown ability to solve many tasks[10], varying from more classical ones like translation and question-answering to more peculiar ones, like common sense reasoning and reading comprehension.

Mathematically, language models seek to maximize the likelihood of seeing some sentence by considering the product of conditional probabilities up to the point of generation

$$p(w_T | \theta) = \prod_{i=1}^T p(w_i | w_{1,..,i-1}, \theta) \quad (3.1)$$

Here we list some of the most popular auto-regressive language models:

- **GPT-3:** GPT-3[11] is an unsupervised AI language model developed by OpenAI. It is the successor to the previous iteration of GPT-2, and when released was the largest language model ever created. GPT-3 first showed that LLMs can be used for few-shot learning across different domains and can achieve state of the art results without building a task-specific model.
- **PaLM:** Pathways Language Model[12] is a 540-billion parameter LM developed by Google. The model was trained through the use of Pathways, a new system that enables highly efficient training of very large neural networks across thousands of TPUs. PaLM few-shot evaluation can outperform or match the finetuned state of the art on a wide array of reasoning tasks.
- **Chinchilla:** Chinchilla[13] is a LM released by DeepMind. The peculiarity of Chinchilla is that it has 'only' 70B parameters, but outperforms

models such as GPT-3 (175B). The authors of the paper demonstrate that for the same computing budget a smaller model trained on more data will perform better.

- **GPT-J**: GPT-J[14] is an open-source alternative to OpenAI’s GPT-3, the model was released by Eleuther AI, a group of independent researchers and has 6B parameters
- **BLOOM**: BLOOM[15] is an open-source auto-regressive LM released by HuggingFace. BLOOM has a stronger focus on languages, differentiating itself from the majority of other LMs that mainly focus on the English language. In total, BLOOM was trained in 46 different languages

Amongst all of these LMs, we picked GPT-J, mainly because it was open-source, free and more lightweight, thus more convenient in case of productive usage. The other main alternative we have tried was GPT-2[9], the predecessor of GPT-3, which is now free to use. However, GPT-J’s generated text was more coherent and more fluent

3.5 GPT-J

The generative model used to create the tickets is GPT-J, an open source 6 billion parameter, autoregressive text generation model trained on The Pile dataset released by EleutherAI.

The Pile dataset[16] is an 825 GiB English text corpus composed by 22 diverse subsets, which can be grouped in 5 categories:

- Academic (*ArXiv*, *PubMed Central*, ...)
- Internet (*Wikipedia*, *StackExchange*, ...)
- Prose (*Bibliotik*, ...)
- Dialogue (*Subtitles*, ...)
- Misc (*Github*, ...)

3.5.1 Parameters

The parameters used for the generation of the next token are:

- *min_length*: minimum number of words created by a round of gpt generation
- *max_length*: maximum number of words created by a round of gpt generation.
- *top_k*: the k most likely next words are filtered and the probability mass is redistributed among only these k words. It helps the model not to go off-topic.
- *top_p*: the next words are sampled from the smallest possible set of words whose cumulative probability exceeds the probability p. Compared to *top_k*, in this case the size of the set of possible next token is not fixed. In practice, usually *top_k* and *top_p* are used together.
- *temperature*: the value T used to module the logits distribution. The higher the value of T , the higher the entropy of the logits distribution will be. In other words, tokens with an high probability will be less probable and tokens with a low probability will be more probable.

$$p_i = \frac{\exp(x_i/T)}{\sum_j \exp(x_j/T)}$$

- *repetition_penalty*: the parameter θ for repetition penalty. 1.0 means no penalty

Given a list of generated tokens G ,

$$p_i = \frac{\exp\left(\frac{x_i}{T \cdot I(i \in G)}\right)}{\sum_j \exp\left(\frac{x_j}{T \cdot I(j \in G)}\right)}$$


$$I(c) = \theta \text{ if } c \text{ is True else } 1$$

So the logits distribution of the token change based on if the token has already been generated before.


- *length_penalty*: *length_penalty* > 0 promotes longer sequences, while *length_penalty* < 0 encourages shorter sequences.
- *no_repeat_ngram_size*: If set to an integer > 0, all ngrams of that size can only occur once
Ex: *no_repeat_ngram_size* = 2

Dear Marie, I would like to have some days off. I would like to... ✗
 Dear Marie, I would like to have some days off. I want... ✓

- *num_beams*: Number of beams for beam search. They beams are the number of 'paths' that are considered when choosing the next token. Even if a token is not the one with the highest probability, it can be chosen as the next token if the complete sentence is considered as more probable than the alternatives.



- *do_sample*: If set to 'False' greedy decoding is used (the most probable token is always chosen). Otherwise, sampling is used (the next token is chosen sampling from the distribution of possible next tokens)



- *bad_words*: List of words that are not allowed to be generated by the model.
- *force_words*: List of words that must be generated in the generation.

The default values assigned to the parameters used are shown in Table 3.2, however the application is set up in order to being able to apply different parameters at each run.

3.5.2 Architecture

The GPT architecture is based on the original Transformers paper. The original architecture introduced two types of transformers blocks: the encoder block and the decoder block.

The encoder block converts an input sequence of tokens into a sequence of embedding vectors, whereas the decoder takes the output of the encoder and generate iteratively an output sequence of tokens.

GPT is pre-trained by predicting the next word based on the previous ones. GPT is decoder-only, which means that is assembled only by a stack of decoder blocks. Each decoder block is composed by:

- **Normalization Layers**: a normalization layer [17] normalize all inputs of a neural network across their features. It has been shown that Layer

Parameter	Value
min length	0
max length	50
top k	50
top p	0.85
repetition penalty	1.2
temperature	1
length penalty	1
no repeat ngram size	0
num beams	1
do sample	True
bad words	[]
force words	[]

Table 3.2: Parameters of GPT-J model

normalization enables smoother gradients, faster training, and better generalization accuracy [18]

x : data sample

d : dimension of data sample

y : output of LayerNorm

ϵ : small number added for stability

$$u = \frac{1}{d} \sum_{i=1}^d x_i$$

$$\sigma^2 = \frac{1}{d} \sum_{i=1}^d (x_i - u)^2$$

$$\hat{x}_i = \frac{x_i - u}{\sqrt{\sigma^2 + \epsilon}}$$

$$y = \gamma \hat{x}_i + \beta$$

where γ and β are parameters that the model learns.

- **Masked Self-Attention Layer:** attention is a mechanism that allows neural networks to assign a different amount of weight to each token in a sequence and process each token as a weighted average of all other

tokens.

In practice three matrixes are calculated:

- Q (Query): the representation of the current token
- K (Key): the representation of all the other tokens, which are matched with the current token
- V (Value): the representation of all the words, used for the weighted-average

The Q , K and V matrixes are initialized as:

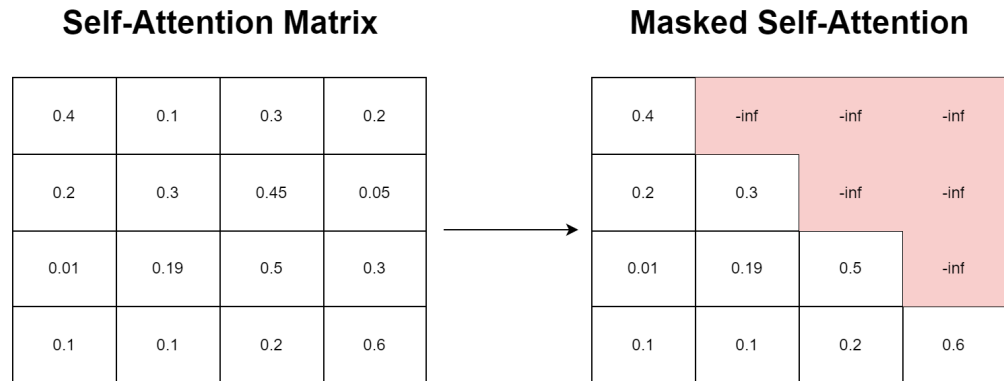
- $Q = W_q X + b_q$
- $K = W_k X + b_k$
- $V = W_v X + b_v$

where X is the input matrix and the other matrixes are randomly initialized and learned by the model. Finally, the attention score is calculated with

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$$

where d is a normalization factor equivalent to the embeddings' dimension. The masked self-attention is a modified version of self-attention where all the tokens that appear after the current one are set to 0, in order not to let the model being influenced by any information regarding the tokens at the next positions. This is fundamental when training generative models such as GPT, whose scope is to predict the successive tokens.

$$Attention(Q, K, V) = softmax\left(\frac{QK^T + Mask}{\sqrt{d}}\right) V$$



- **Feed Forward Neural Network Layer:** Used to add non-linearity to the transformer block

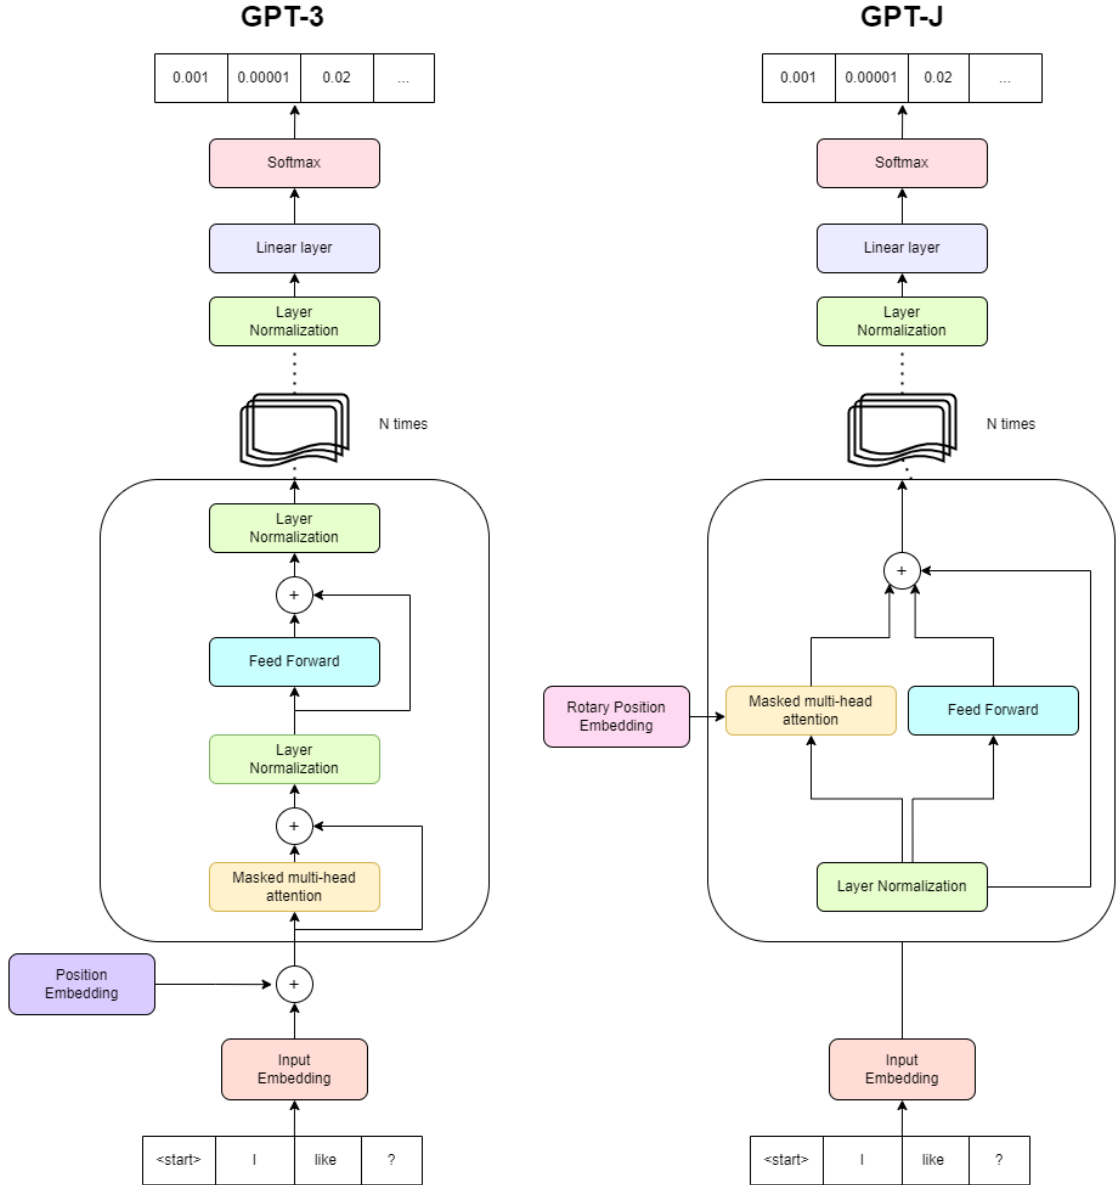


Figure 3.5: GPT-3 and GPT-J architectures compared

GPT models use a Byte-Pair Encoding tokenization. The Byte-Pair Encoding algorithm starts by building a vocabulary with all the single characters of the corpus we are training on. Then, at each step of the

Example: the corpus is formed by the words book, nook, noob, boob and books respectively 12,8,14,5 and 6 times
The initial vocabulary will be
{'b', 'o', 'k', 'n', 's'}

Word	Characters	Counts
book	b o o k	12
nook	n o o k	8
noob	n o o b	14
boob	b o o b	5
books	b o o k s	6

The character combination that appears the most times consecutively is 'oo', so our vocabulary becomes
{'b', 'o', 'k', 'n', 's', 'oo'}




Word	Characters	Counts
book	b oo k	12
nook	n oo k	8
noob	n oo b	14
boob	b oo b	5
books	b oo k s	6

Now the character combination that appears the most times is created by unifying 'oo' and 'k'
Updated vocabulary:
{'b', 'o', 'k', 'n', 's', 'oo', 'ook'}



And so on, until we reach the vocabulary length desired



Word	Characters	Counts
book	b ook	12
nook	n ook	8
noob	n oo b	14
boob	b oo b	5
books	b ook s	6

Figure 3.6: Byte-Pair Encoding Tokenization Example

algorithm, the two tokens which appear consecutively the most in the words of the corpus are unified and create a new token. This process continue until the desired vocabulary length is satisfied. An example taken from the site [19] is shown in the Figure 3.6

Compared to GPT-3, GPT-J[14] has two minor architectural differences (shown in Figure 3.5):

- Rotary Embedding
- The attention layer and the feedforward layer in parallel for decreased

communication

Rotary Position Embedding

Position embeddings are used to infer the notion of position to the model, which does not have any sense of position for each token. In other words, using the attention mechanism each token "match" with the other tokens in the same manner, not considering where if the other token is located right after the current one or if it is at the end of the sentence. Position embeddings are used to add to the model this sense of position.

Rotary position embedding has been introduced by Su et al.[20], it is a novel method that unifies absolute and relative approaches to position embeddings. The typical approach, which is also used by GPT-3, is to use a sinusoidal embedding, which is defined as

$$\begin{cases} p_{1,2t} = \sin(k/10000^{2t/d}) \\ p_{1,2t+1} = \cos(k/10000^{2t/d}) \end{cases}$$

where $p_{1,2t}$ is the 2^{th} element of the d -dimensional vector p_i RoPE instead proposes to incorporate the relative position information by multiplying the context representation with the sinusoidal functions instead of directly adding them.

If we define

$$\begin{aligned} q_m &= f_q(x_m, m) \\ k_n &= f_k(x_n, n) \end{aligned}$$

where f_q and f_k are functions that incorporate the m^{th} and n^{th} positions respectively to the vector embeddings x_m and x_n to produce the query and key vectors, we can require the inner product of the query q_m and k_n to be formulated by a function g that depends only on the word embeddings x_m , x_n and their relative position $m - n$

$$\langle f_q(x_m, m), f_k(x_n, n) \rangle = g(x_m, x_n, m - n)$$

In the simplest case $d = 2$, $f_{\{q,k\}}$ are defined as

$$f_{(\{q,k\},\{m,n\})} = (W_{\{q,k\}} x_{\{m,n\}}) e^{i\{m,n\}\theta}$$

and therefore we obtain

$$g(x_m, x_n, m - n) = \text{Re}[(W_q x_m)(W_k x_n)^T e^{i(m-n)\theta}]$$

which preserves the relative positional information of the word embeddings. This equation is used when calculating the self-attention, which will become

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{g(x_m, x_n, m - n)}{\sqrt{d}}\right) V$$

This equation can be generalized for $d > 2$, as shown in the original paper. In the end, incorporating the RoTE is pretty straightforward, you just have to rotate the word embedding by a multiple of its position index. According to the researchers that published GPT-J[21], using RoTE leads to a faster convergence of training and validation losses and a lower overall validation loss.




Figure 3.7: Implementation of RoPE, Image taken from original paper[20]

3.6 Architecture analysis

To understand better how the model was behaving and why it gave certain types of outputs rather than others, we used the python library Ecco[22],

which creates interactive visualizations that show at which layer of the architecture the final token has been decided, which input tokens contribute the most for a prediction, and many other insights on the model.

In particular, we used two different methods:

- **Input Saliency:** used to show how much did each input token contribute to producing the output token
- **Neuron Activation Analysis:** used to examine underlying patterns in neuron activations using non-negative matrix factorization

Input Saliency

To get a better grasp of the most useful information of the prompt, and to understand how we could modify it to achieve better results, we used the Gradient * input[23] method. This technique calculates the partial derivatives of the output of the model and multiplies them with the input itself. Then the inputs with the highest scores are considered the ones that influenced the most the generation of the new tokens

$$score = x_i \nabla f(x_i)$$

where f is the architecture output.

The main problem with the Gradient * input method is that only one input is considered. Integrated Gradients[24] solve this issue, computing the average gradient while the input varies along a linear path.

$$score = (x_i - x'_i) \int_{\alpha=0}^1 \nabla f(x' + \alpha(x - x')) d\alpha$$

Nevertheless, we used the Gradient * input method for computational issues (The Integrated Gradients took too much RAM of the GPUs).

In the images under, we underline some of the considerations we have done while defining the prompts.

```

From:vbarbe@weber.fr\n
To:hr@HumbertSA.com\n
Firstname:Victoire\n
Lastname:Barbe\n
Company:HumbertSAFrance\n
Date:13/04/2019\n
Ticketcategory:Complaint\n
Ticketsub-category:complaint\n
Complaint:Complaintaboutasuperior\n
Subject:mybossismakingmeworkwaybeyondtheworkinghours\n
DearSir/Madame,I don't want to work anymore with >> myboss.\n

```

Figure 3.8: Input Saliency: First token

To create the first token, the model focus on the elements that resemble the content of an email, in particular on the fact that it is indeed a 'Ticket'

```

From:alyssa.blankenship@owens.info\n
To:hr@StewartGroup.com\n
Firstname:Alyssa\n
Lastname:Blankenship\n
Company:StewartGroupUnitedStates\n
Date:21/07/2021\n
Ticketcategory:Lifevent\n
Ticketsub-category:Healthissues\n
Datestartabsence:27/07/2021\n
Reasonabsence:amedicalconsultation\n
Subject:Requestforsickleavefor8days\n
DearSir/Madame, mynameisAlyssaBlankenshipandIworkatStewartGroupUnitedStates.Iamrequesting >> a  

leaveofabsenceon8daysfrom27July2020.IsisthenormalprocedureanddoIneedtogetapprovalfromtheHRor  

theHRisnotallowedtoextendFMLAleavetime.IworkintheHR,Ionlyget

```

Figure 3.9: Input Saliency: Subject

Clearly stating the subject of the ticket helps the model create tokens that are on the right topic (In this example 'days of absence')

Neuron Activation Analysis

The Feed Forward Neural Network layer is one of the major components inside a transformer's block. To better understand how the neurons of different layers were 'activated' and how the neurons contributed towards each generated token, we exploited the Factor Analysis provided by Ecco.

Firstly, Ecco calculates the activation scores of each neuron over all layers, and then uses Non-negative Matrix Factorization (Figure 3.11) to do dimensionality reduction on the matrix of activations, which will be reduced to a

```

From: alyssa.blankenship@owens.info\n
To: hr@StewartGroup.com\n
Firstname: Alyssa\n
Lastname: Blankenship\n
Company: StewartGroupUnitedStates\n
Date: 21/07/2021\n
Ticketcategory: Lifeevent\n
Ticketsub-category: Healthissues\n
Datestartabsence: 27/07/2021\n
Reasonabsence: amedicalconsultation\n
Subject: Request for sick leave for 8 days\n
Dear Sir/Madame, my name is Alyssa Blankenship and I work at StewartGroupUnitedStates. I am requesting >> a
leave of absence on 8 days from 27 July 2020. Is this the normal procedure and do I need to get approval from the HR or
the HR is not allowed to extend FMLA leave time. I work in the HR, I only get

```

Figure 3.10: Input Saliency: HR

Clearly stating that we are writing to hr can help the model know the context, and consecutively use a certain language, specific words...

matrix $M \times T$, where M is a parameter we decide and T is the number of tokens, starting from a matrix $(n \cdot N) \times T$, where n is the number of neurons per layer and N is the number of layers. In GPT-J $n = 16384$ and $N = 28$.

Then, for each factor, which are the dimensionality-reduced layers, we

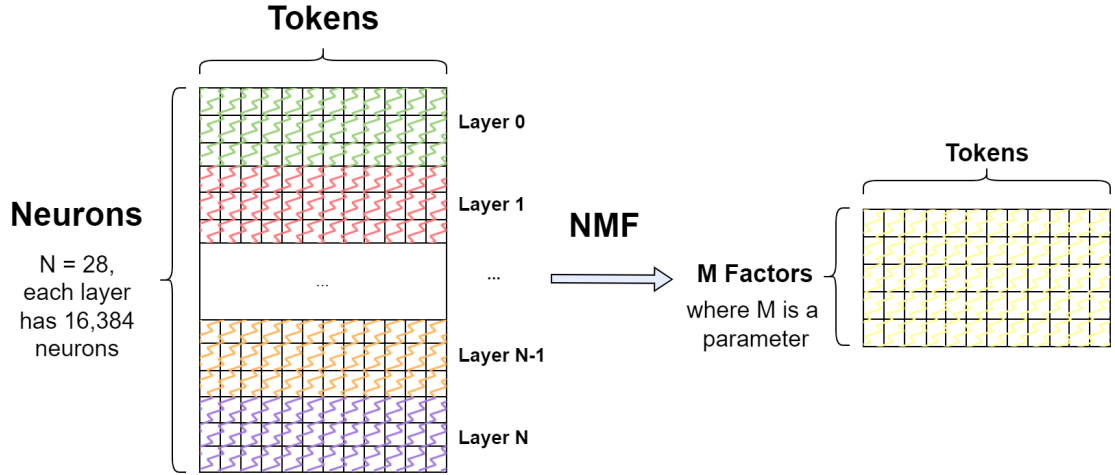


Figure 3.11: Non-negative Matrix Factorization on Activation Matrix

visualize their activated neurons. From the example in the Figure 3.12, we can show for each factor what their main contributions are:

1. Punctuation
2. Mail prompt (From, To, First Name, Last Name, Company...)
3. New lines
4. Dates
5. First token, common across various GPT factors
6. Medical terms (usually terms related to the tickets' topic)
7. People's names and contacts
8. Company's name and contacts
9. Ticket's subject (Category, Sub-category and additional info)
10. Email presentation ('Dear Sir/Madame...')

From : a lyss a . blank ens hip @ ow ens . info | To : hr @ Stew art Group . com | First name : A lyss a | Last name : Blank ens
hip | Company : Stewart Group United States | Date : 21 / 07 / 20 21 | Ticket category : Life event | Ticket sub - category :
Health issues | Date start absence : 27 / 07 / 20 21 | Reason absence : a medical consultation | Subject : Request for sick
leave for 8 days | Dear Sir / Mad ame , my name is A lyss Blank ens hip and I work at Stewart Group United States . I am
requesting >> leave for 8 days , starting 27 th of

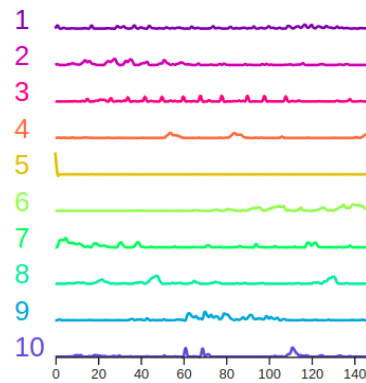


Figure 3.12: Neuron Activation Analysis

3.7 Survey

In order to understand whether the tickets were similar to real tickets, we conducted a survey internal to the company. We were not able to get real data from the company’s HR due to the sensitive nature of the data. In order to avoid any ethical implications, we decided to conduct our assessment without involving any real personal data or situations.

We asked to some to act as the automatic text generator, filling out a form in an Excel file, where the users were given twenty random prompts, similar to the ones that are given to the GPT model.

After giving a brief explanation of the project and what we were trying to achieve, we recommended the users not to use their personal information while writing the fake tickets, but to use the same style of text and the same vocabulary they would use in a real situation.

For each ticket, there were several information attached that appeared over all the tickets, which were: name of the person, category of the ticket, sub category of the ticket and some information that were specific to the ticket’s category, for example for the category-sub_category ‘Salary - Salary Raise’ the additional information given was the new requested salary.

We asked the users not to use necessarily all the information that were present in the prompts, but only the ones that felt natural to use, and to slightly change them if it made sense to them (Ex: Sick leave from Oct 20 2022 can become ‘from this Thursday, the 20’)

Each user was given twenty different prompts sampled randomly from a set of 800 total prompts, belonging to the same categories-sub_categories of our taxonomy. The prompts contained only the contextual information of the ticket and no ticket’s text, so, unlike the prompts used for the GPT models, there were no initial text such as ‘Dear Sir/Madame, I would like to ...’

The respondent were also asked to tick the information that they used, and these information were supposed to be used as testing in the NER use case. However, many people either did not tick any information or they ticked only partially what they used, therefore in the end we did not make use of these data.

Figure 3.13: Example of survey

3.8 Results

In this section we analyze the created dataset, comparing it to the tickets we collected from the survey.

The final dataset of HR tickets automatically generated is composed by exactly 16,000 tickets, 2000 for each category. On the other hand, we collected 259 test tickets from 29 different respondents.

To evaluate the generation of the tickets, we have used a handful of traditional unsupervised text evaluation techniques. We were not able to exploit more complex classical metrics, such as BLEU and ROUGE, not having real data as references for each prompt. The reported metrics are:

- avg ttr unigram: average TTR(type token ratio) for the unigrams of the tickets

$$TTR = \frac{\text{number_of_unique_unigrams}}{\text{total_number_of_unigrams}}$$

- avg ttr bigram: average TTR(type token ratio) for the bigram of the tickets

$$TTR = \frac{\text{number_of_unique_bigrams}}{\text{total_number_of_bigrams}}$$

- avg ratios of nouns: average ratio of nouns in the tickets

$$\text{noun_ratio} = \frac{\text{counter_noun_words}}{\text{counter_all_words}}$$

- avg ratios of verbs: average ratio of verbs in the tickets

$$\text{verb_ratio} = \frac{\text{counter_verb_words}}{\text{counter_all_words}}$$

- avg_word_frequency: average frequency of words, using as frequency estimates a dump of the English version of Wikipedia, considering only words appearing at least 10 times

$$\text{word_freq} = \begin{cases} \log_e(\text{word_freq_wiki}), & \text{if } \text{word_freq_wiki} \geq 10 \\ \text{skip}, & \text{otherwise} \end{cases}$$

First, we reported the type-token-ratios given that high-quality writing has been associated with the presence of more diverse words and phrases[25]. We computed them for each ticket and then we averaged the results.

Second, since lower frequency words indicate a more advanced output[26], we compute the average word frequency of the generated words.

Then, we calculate nouns and verb ratios over sentence length, as indicators of syntactic complexity, and therefore richer text[27]. To single out nouns and verbs, we use the pre-trained Spacy POS model.

We report in Table 3.3 the results we obtained initially. The metrics have been calculated not only on the tickets generated by the model and the tickets of the surveys, but also on some baseline datasets (*Amazon reviews*, *Reddit comments* and *NIPS papers*). The datasets have been chosen to be as diversified as possible in terms of text longevity, type of language and terms used.

The main takeaways from the metrics are:

- The model tends to write more unique words and not to repeat itself, since the average TTR of both unigrams and bigrams are higher in the tickets of the survey
- The model tends to use more nouns than a real person
- The model tends the same number of verbs as a real person
- The average number of words (word count) varies a lot from category to category, both in generated and survey tickets

We know that these analyses are not completely reliable due to the low number of real data compared to the size of our dataset. However they are a good indication of how to improve the generation.

For this reason, we changed the parameters of the ticket creation and recreated the dataset from scratch, in order to be more akin to the survey tickets.

The final results are shown in Table 3.4.

Some of the surveys' metrics change a bit compared to the first table, this is due to some late respondents to the survey, which skewed a little bit the final test metrics.

Topic analysis

Topic modeling is a type of statistical modeling technique used to discover the hidden topics and patterns in a corpus of text. It is a type of unsupervised machine learning method that automatically discovers topics from a collection of documents by analyzing words and phrases in each document. Topics are usually represented by word clusters or distributions that are related to each other. Topic modeling can be used to generate meaningful summaries about large collections of documents.

We used topic modeling to visualize in a summarized form the major topics of our dataset and to verify that the topics created correspond to the ones in the taxonomy.

The model of topic modeling we chose is BERTopic[28]. BERTopic generates document embedding with pre-trained transformer-based language models, clusters these embeddings, and finally, generates topic representations with the class-based Term Frequency-Inverse Document Frequency (TF-IDF) procedure

BERTopic

- 1: Documents are embedded to create representations in vector space using Sentence-BERT
 - 2: The embeddings generated are reduced in dimensionality using UMAP
 - 3: The reduced embeddings are clustering used HDBSCAN. HDBSCAN is a variation of the hierarchical clustering algorithm DBSCAN that models clusters using a soft-clustering approach, allowing noise to be modeled as outliers
 - 4: The most relevant terms for each topic are found using a class-based version of TF-IDF. All documents in a cluster are treated as a single document by simply concatenating them, so we measure how much information a term provides to a class instead of a document.
 - 5: If the user requests a number of topics smaller than the number of topics generated, the least common topics are iteratively merged with their most similar topic, until the number of topics requested is satisfied.
-

Dataset	avg ttr unigram	avg ttr bigram	avg ratios of nouns	avg ratios of verbs	avg word frequency	avg word count
Amazon reviews	0.81	0.98	0.16	0.10	13.86, $\sigma=0.47$	73.73, $\sigma=71.75$
Reddit comments	0.91	0.99	0.15	0.11	13.33, $\sigma=1.53$	37.70, $\sigma=84.96$
NIPS papers	0.26	0.72	0.18	0.06	13.50, $\sigma=0.25$	5798.40, $\sigma=737.17$
Tickets survey	0.87	0.99	0.17	0.11	13.84	42.73, $\sigma=28.83$
Tickets generated by GPT	0.93	1.00	0.14	0.10	13.84, $\sigma=0.35$	57.87, $\sigma=19.13$
Survey Salary_Salary raise	0.85	0.99	0.20	0.10	13.85, $\sigma=0.56$	55.65, $\sigma=35.67$
GPT Salary_Salary raise	0.90	1.00	0.15	0.10	13.72, $\sigma=0.33$	70.73, $\sigma=16.36$
Survey Complaint_Complaint	0.83	0.98	0.15	0.13	13.66, $\sigma=0.91$	63.64, $\sigma=46.67$
GPT Complaint_Complaint	0.91	0.99	0.14	0.12	13.94, $\sigma=0.32$	72.81, $\sigma=21.62$
Survey Life event_Personal issues	0.89	0.99	0.17	0.12	13.99, $\sigma=0.56$	33.54, $\sigma=19.85$
GPT Life event_Personal issues	0.92	1.00	0.13	0.12	14.08, $\sigma=0.25$	63.88, $\sigma=15.62$
Survey Refund_Refund travel	0.90	0.99	0.17	0.11	13.56, $\sigma=0.56$	34.76, $\sigma=17.02$
GPT Refund_Refund travel	0.96	1.00	0.16	0.08	13.72, $\sigma=0.38$	38.91, $\sigma=10.83$
Survey Timetable change_Shift change	0.83	0.99	0.16	0.10	14.16, $\sigma=0.66$	38.15, $\sigma=15.30$
GPT Timetable change_Shift change	0.93	1.00	0.13	0.10	13.79, $\sigma=0.37$	47.95, $\sigma=9.72$
Survey Ask information_Accommodation	0.88	0.99	0.17	0.11	13.80, $\sigma=0.50$	39.93, $\sigma=23.89$
GPT Ask information_Accommodation	0.95	1.00	0.14	0.10	13.68, $\sigma=0.36$	47.10, $\sigma=10.59$
Survey Life event_Health issues	0.90	0.99	0.17	0.13	13.88, $\sigma=0.44$	31.92, $\sigma=14.70$
GPT Life event_Health issues	0.91	1.00	0.14	0.10	13.84, $\sigma=0.33$	59.07, $\sigma=16.06$
Survey Salary_Gender pay gap	0.85	0.97	0.21	0.11	13.87, $\sigma=0.43$	53.94, $\sigma=32.71$
GPT Salary_Gender pay gap	0.93	0.99	0.15	0.11	13.96, $\sigma=0.28$	62.52, $\sigma=18.67$

Table 3.3: Initial Results

The formula of the class-based TF-IDF is:

$$W_{t,c} = tf_{t,c} \cdot \log(1 + \frac{A}{tf_t})$$

where $tf_{t,c}$ is the frequency of term t in class c , A is the average number of words per class and tf_t is the frequency of term t across all classes.

Dataset	avg ttr unigram	avg ttr bigram	avg ratios of nouns	avg ratios of verbs	avg word frequency	avg word count
Amazon reviews	0.81	0.98	0.16	0.10	13.86, $\sigma=0.47$	73.73, $\sigma=71.75$
Reddit comments	0.91	0.99	0.15	0.11	13.33, $\sigma=1.53$	37.70, $\sigma=84.96$
Tickets survey	0.86	0.99	0.17	0.11	13.89, $\sigma=0.59$	44.43, $\sigma=27.46$
Tickets generated by GPT	0.94	1.00	0.14	0.10	13.86, $\sigma=0.39$	49.22, $\sigma=15.49$
Survey Salary_Salary raise	0.84	0.99	0.21	0.10	13.89, $\sigma=0.52$	55.10, $\sigma=32.44$
GPT Salary_Salary raise	0.90	1.00	0.15	0.10	13.73, $\sigma=0.35$	62.87, $\sigma=14.12$
Survey Complaint_Complaint	0.82	0.98	0.15	0.13	13.74, $\sigma=0.84$	62.74, $\sigma=42.83$
GPT Complaint_Complaint	0.92	0.99	0.14	0.11	13.94, $\sigma=0.33$	59.46, $\sigma=18.39$
Survey Life event_Personal issues	0.88	0.99	0.17	0.12	14.03, $\sigma=0.55$	38.15, $\sigma=20.86$
GPT Life event_Personal issues	0.94	1.00	0.14	0.11	14.19, $\sigma=0.30$	45.49, $\sigma=16.94$
Survey Refund_Refund travel	0.88	0.99	0.17	0.11	13.61, $\sigma=0.53$	37.50, $\sigma=17.36$
GPT Refund_Refund travel	0.96	1.00	0.16	0.08	13.72, $\sigma=0.38$	39.52, $\sigma=10.46$
Survey Timetable change_Shift change	0.83	0.99	0.16	0.10	14.17, $\sigma=0.62$	38.73, $\sigma=14.71$
GPT Timetable change_Shift change	0.93	1.00	0.12	0.10	13.74, $\sigma=0.40$	42.80, $\sigma=7.55$
Survey Ask information_Accommodation	0.87	0.99	0.16	0.11	13.90, $\sigma=0.50$	41.86, $\sigma=23.79$
GPT Ask information_Accommodation	0.95	1.00	0.14	0.10	13.65, $\sigma=0.38$	46.79, $\sigma=10.99$
Survey Life event_Health issues	0.87	0.99	0.17	0.13	13.94, $\sigma=0.42$	37.91, $\sigma=21.80$
GPT Life event_Health issues	0.94	1.00	0.16	0.09	13.95, $\sigma=0.37$	43.36, $\sigma=11.26$
Survey Salary_Gender pay gap	0.85	0.97	0.22	0.11	13.87, $\sigma=0.41$	53.14, $\sigma=31.13$
GPT Salary_Gender pay gap	0.94	1.00	0.14	0.11	13.99, $\sigma=0.29$	53.49, $\sigma=13.51$

Table 3.4: Final Results

In Figure 3.14 we show all the topics that BERTopic outputs, and in Figure 3.15 the grouped major 8 topics.

We are satisfied with the results, we can clearly recognize 7 of our 8 categories in the topics created, the only one missing is the request of time off due to personal reasons, that have been incorporated in Topic 3.

The topics generated by BERTopic are:

- **Topic 0:** Request of shift change
- **Topic 1:** Request of explanation of wage gap between genders
- **Topic 2:** Official complaint about colleague/superior
- **Topic 3:** Request of time off due to health reasons
- **Topic 4:** Request of salary increase
- **Topic 5:** Request of refund for work travel
- **Topic 6:** Request of refund for work travel
- **Topic 7:** Request of information for new accommodation




Figure 3.14: Hierarchy of all topics generated by BERTopic

3.9 Additional Experiments

In this section, we are going to describe various methods we have tried for the generation of tickets which have not been included in the final version of the ticket generation for different reasons.

Some of these reasons are lack of computational power, not good enough results and lack of time.




Figure 3.15: Main 8 topics generated by BERTopic

3.9.1 Topical Generation

Zandie et al.[29] have presented topical language generation, which consists of generating text conditioned on a specific chosen topic. The topic can be chosen by the users from a predetermined list of topics. The topics are represented as a distribution over a vocabulary, and they are extracted from a corpus using algorithms of Topic Modeling.

The algorithm used is Latent Semantic Analysis, which gives us scores for each token per topic. LSA is able to discover latent topics in a document-term matrix by utilizing Singular Value Decomposition (SVD) to decompose the matrix.

The topics are inferred to the text generation by adding the embedding of the topic to the logits of the model before the last softmax layer, where the parameter γ control the topic inference (the higher the parameter γ is, the more the text generated will be about the topic chosen)

$P(t_j | x_i)$: probability of topic j given word k

$S(x_i | x_{<i})$: logits of token i

They used a threshold to increase the probability of only the tokens that

already have a minimum probability to be sampled.

$$\logprob(i) = \begin{cases} \log P(t_j | x_i), & \text{if } S(x_i | x < i) \geq \text{threshold} \\ 0, & \text{otherwise} \end{cases}$$

The final equation is:

$$P(x_i | x < i, t_j) = \text{softmax}(S(x_i | x < i) + \gamma \logprob(i))$$

We experimented this method to generate the tickets' texts. The topic-word matrix was calculated starting from 20 Newsgroups, a dataset of 18000 news posts on 20 topics.

Topic used: football

[...]

Ticket category: Life event

Ticket sub-category: Health issues

Date start absence: 24/08/2016

Reason absence: a physiotherapy visit

Subject: Request for sick leave for 1 day

Dear Sir/Madame, my name is Adriana Giulietti. I am requesting football medical treatment on the following dates : - 28 nfl season 2016 (from 30th of August regular football game) professional training session in Milan from 29 football football league match between FC Portcullis football conference matches and other

Topic used: president, government

[...]

Ticket category: Life event

Ticket sub-category: Health issues

Date start absence: 25/12/2012

Reason absence: Hordeolum

Subject: Request for sick leave for 2 days

Dear Sir/Madame, my name is Cristian Manso. I am requesting sick pay or vacation from you because state government has decided to cancel the contract which was signed on 20 th of December 2012 with head office called president and deputy leader states house republic functions as a general administration department in charge common services like public works system

The possible future implementations of the ticket generation could be completely executed with a unique prompt, letting the user choose the topic from a set of pre-trained topics. This would be more flexible than the current method since to add new categories it would be sufficient to add a new row to the topic-word matrix.

However, the current results were not satisfying: the created tickets were not fluent and we did not find a corpus to create a topic-word matrix of topics related to the HR ticketing environment.

3.9.2 Fine-tune GPT

GPT can be adapted to a particular task by fine-tuning, either by supplying it with a dataset tailored to said task, or by manually adjusting the model's parameters. This process allows the model to be customized and to achieve much better results on the desired task.

As mentioned before, we want to generate tickets with a language and style similar to emails. Fine-tuning the model on the *Enron dataset* should improve the performance of the model at generating new Enron-style emails compared to the generic trained model.

We tried fine-tuning GPT-J on only the *_sent_mail* folders of the dataset. However, the tickets generated were not much different than the ones created by the base model. The evaluation was carried out only by looking at the tickets generated one by one, without using any sort of metrics.

After various tests, we ascertained that with prompt tuning we could have equal results as with the fine-tuning of the model, avoiding to spend time to fine-tune the model.

Chapter 4

Use Cases

In order to assess the suitability of the HR ticket dataset for downstream machine learning applications in HR, we exploit the obtained dataset using as labels the categories of the tickets.

The goal of testing different downstream machine learning tasks on the dataset is to show that our dataset is sufficiently rich to enable meaningful learning.

All the experiments have been carried out using as training data the dataset generated by the model, whereas the test data used was the tickets collected with the survey.

One of the main motivations to acquire the survey tickets was to have data that had no biases, or at least biases different from ours. Using a portion of the HR tickets dataset as test data would be meaningless, since the model would not necessarily be tested with data that was outside of the one it had been trained with. On the contrary, achieving good results on a portion of the HR ticket dataset would assert the performance of the models used and not the usefulness of the dataset.

4.1 Classification

Sentence classification is a natural language processing task that involves assigning a predefined category or label to a given sentence. One way to approach sentence classification is to use word embeddings, which are numerical representations of words that capture their meaning and context within a sentence. Word embeddings can be generated using various techniques, such as training a neural network on a large dataset of text or using a pre-trained language model.

Once the word embeddings have been generated, you can input them into a classifier along with other features of the sentence, such as its grammatical structure, to make predictions about the sentence's category or label. The classifier uses the word embeddings as input to make predictions about the sentence's category or label.

We used two different pre-trained language models: BERT and fastText. We chose BERT because it is a widely used and well-studied model, and our objective was not to find the absolute best classifier, but only to show that the dataset could be used as training data.

Instead, we chose fastText to have an option that was CPU-friendly, and that could be trained in a matter of a few minutes.

The training data used were the texts of the 16000 tickets that compose the HR ticket dataset, whereas the test data were the survey tickets. The labels were the combination of ticket category and ticket subcategory of each ticket (Ex. "Life event_Health issues").

4.1.1 FastText Classifier

FastText is a library developed by Facebook AI Research that provides a set of training algorithms for supervised learning of word embeddings from raw text and also can be used to build and train supervised text classification models.

FastText creates word embeddings using a combination of character n-grams and word n-grams. To do this, fastText first breaks each word in the training text into its constituent character n-grams. For example, the word "cat" might be broken into the character n-grams "c", "ca", "cat", "a", "at", and "t". These character n-grams are then used to generate vectors. The word

embeddings are calculated as the sum of their n-gram vectors.

To learn the embeddings of these n-grams, fastText uses a skip-gram model, which is a type of model that maximizes the probability of observing a word given its context, or in other words its surrounding words. By using this approach, fastText is able to learn high-quality word embeddings that capture the syntactic and semantic information of words.

The fastText classification uses the fastText embeddings as input to the model, and in particular:

fastText Classification

- 1: Words representation are averaged into a text representation
 - 2: The text representation is fed to a linear classifier
 - 3: The output of the linear classifier is given as input to a softmax function to compute the probability distribution over a set of predefined classes
-

4.1.2 BERT classifier

BERT (Bidirectional Encoder Representation from Transformers) is a model developed by Google AI based on the transformers architecture. BERT was trained on a large corpus on two different tasks: Masked Language Modeling and Next Sentence Prediction.

Through the use of MLM, in the training phase, 15% of the tokens in a sentence are obscured and BERT can then be utilized to utilize the surrounding words in both directions to predict the masked tokens, facilitating bidirectional learning from the text.

On the other hand, NSP helps the model understand the relationships between sentences. Specifically, in the training phase, 50% of the examples are actually successive sentences, while in the other 50% of the cases this condition is not respected.

BERT is quite straightforward to fine-tune, we just need to plug in the necessary layers at the top of the BERT architecture and fine-tune for a sufficient number of epochs the model end-to-end.

The first token of a BERT embedding is the [CLS] token, which is a special token that acts as an aggregate representation of the sentence. The final embedding of the [CLS] token is used as input for the classification task.

The most simple method, which is also the method used by the BERT-ForSequenceClassification class by HuggingFace, is to take the embedding of the [CLS] token at the last hidden layer and fed it to a single layer of a feed-forward network. The layer of the feed-forward network will have n input, where n is the dimension of the embedding of a token, and m output, where m is the desired number of outputs.

The final model is shown in Figure 4.1




Figure 4.1: Schema of BERT classification

4.1.3 Results

In this section we show the results of the classification experiments. We want to reiterate that the point of the experiments was not to get the best

possible results, therefore the models were not trained for a large number of epochs and we did not try an excessive number of possible hyperparameters. The hyperparameters tried are shown in Table 4.1 and Table 4.2. If a hyperparameter is not shown in the table it means we have used its default value.

The technique used for finding the optimal hyperparameters is Grid Search, which means we have trained and tested the models for each combination of all the chosen values of the hyperparameters to determine the hyperparameters' set which achieves the best results.

Size of context window (ws)	5
epochs	20
minimal number of word occurrences (minCount)	1
max length of word ngram (wordNgrams)	3
learning rate (lr)	0.5
learning update rate (lrUpdateRate)	100
sampling threshold (t)	0.0001

Table 4.1: Hyperparameters fastText

epochs	[3, 5]
learning rate	[1e-05, 5e-04, 1e-04, 5e-05, 5e-06]
weight decay	[0.001, 0.01, 0.05]
train batch size	[8, 16]
warmup steps	500

Table 4.2: Hyperparameters BERT

As we expected, fastText performs worse than BERT, achieving an f_1 score of only 0.41.

On the other hand, the BERT classifier with hyperparameters

- epochs: 5
- learning rate: 5e-05
- weight decay: 0.001
- train batch size: 8

- warmup steps: 500

achieved an f_1 score of 0.78. We show also the confusion matrix(Figure 4.2) to highlight how the majority of errors come from two sub-categories that belong to the same category, in particular, "Gender wage gap"- "Salary raise" and "Personal issues"- "Health issues". In some way, this validates our initial decision on the taxonomy of the tickets.




Figure 4.2: Confusion matrix of Ticket classification with BERT

4.2 Anonymization

Anonymization of personal data refers to the process of removing personally identifiable information (PII) from data sets so that the individuals represented in the data cannot be identified. This is accomplished by either completely removing or replacing identifiable data with generic values. The purpose of anonymization is to protect the privacy of individuals while still allowing the data to be used for legitimate purposes.

Most of the recent implementations of anonymization work by masking the personal information of a person, such as names and surnames, telephone numbers, addresses, credit card numbers and so on.

One of the most famous libraries for anonymization is Presidio, developed by Microsoft. Presidio exploits pattern recognition with regex and Named Entity Recognition to find all the personal information and mask them.

The main disadvantage of such techniques is that often the personal subject can be identified through the so-called quasi-identifiers, that more often than not are not masked.

Here reported some examples that are not masked by Presidio:

Original sentence

The new intern at my office, the one with red hair, caught covid last week

Sentence redacted by Presidio

The new intern at my office, the one with red hair, caught covid
<DATE_TIME >

Original sentence

The boss of the HR department has made some weird comments about how I dress

Sentence redacted by Presidio

The boss of the HR department has made some weird comments about how I dress

Our new approach exploits a Sequence to Sequence model called T5, which stands for Text-To-Text Transfer Transformer. T5 is a model released by Google in 2019, it is a standard encoder-decoder transformer that, unlike BERT, always returns strings as outputs. This is why it is called a sequence-to-sequence model.

T5 is a unified model that can be applied for many downstream tasks, such as sentiment analysis, sentence completion, question answering. . .

In T5 architecture there are *adapted layers* after each feed-forward layer, whose scope is to diminish the number of parameter updates for each fine-tuned model. In fact, *adapted layers* are dense-ReLU-dense blocks that are designed so that their input dimensionality and output dimensionality are equal. This lets us insert them into the Transformer architecture without other changes. When fine-tuning for each different task, only the *adapted layers* and the normalization layers will be updated. This results in a considerable reduction of parameter updates.

We followed a few-shot learning approach with T5. Rather than relying on a large amount of training data to allow a pre-trained model to adjust to a given task accurately, few-shot learning employs the use of a few examples to direct the training of a machine learning model with very minimal data. For each category of tickets, we wrote 10 examples of "anonymized" tickets, where we removed all personal information and information that could be retraced to the original writer, maintaining only information that could be useful for analysis purposes. Here are a couple of examples:

Original sentence:

Dear Sir/Madame, I cannot stand anymore this discrimination and prejudice against French people at work

Anonymization used for few-shot learning:

The employee is filling a complaint for discrimination based on nationality

Original sentence:

Hello, my name is Zacaredas Pinilla and I work at Laguna-Franco Spain. I am having trouble finding accommodation in the Algeciras area so if you could help me with this matter it would be greatly appreciated

Anonymization used for few-shot learning:

The employee is asking for help finding an accommodation

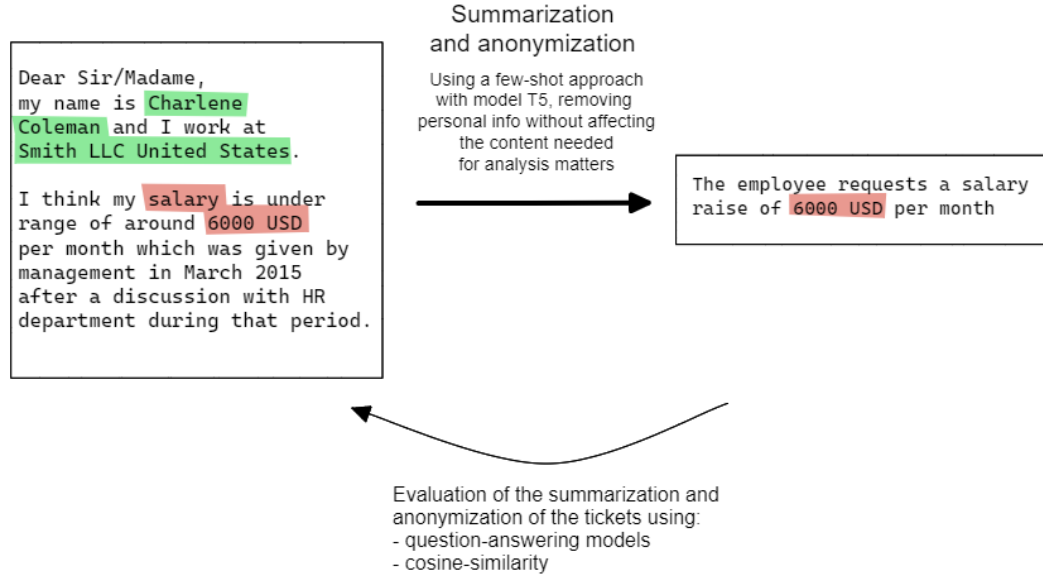


Figure 4.3: Schema of Ticket Anonymization

To evaluate the goodness of anonymization and how much of the information has remained after the anonymization we have tested three different methods:

- **QAGS:** Questions Answering and Generation for summarization
- **SUMMAQA:** Unsupervised metric for reinforced summarization model
- **Cosine Similarity:** Cosine similarity of embeddings

These models were created originally to evaluate the goodness of a summarization, we have adapted them to be used for evaluating our anonymization. The first two methods follow the same philosophy: they generate questions to ask both the original version and the anonymized version of the tickets, looking if the answers are coherent. We call them Questions and Answers models. However, they have a few key differences: SummaQA does not create natural sentence questions but uses as questions the masked version of the sentences. Instead, QAGS creates natural language questions using a pre-trained model.

4.2.1 QAGS

QAGS is an automatic evaluation protocol that is based on the assumption that if we reply to questions considering as factual bases the summary and the source, we can determine if the summary/anonymization is consistent with its source by comparing the answers.

QAGS works like this:

QAGS

- 1: A BART[30] model is fine-tuned for question generation: the model receives both the answers and the source article from the NewsQA dataset and is trained to maximize the likelihood of the paired question
 - 2: At test time, named entities and noun phrases are extracted from the context using spaCy and are considered as answers candidates. The summary is used as the context.
 - 3: The BART model generates questions based on the summary and its entities
 - 4: A BERT model is fine-tuned for Question Answering on the SQuAD dataset
 - 5: We answer with the help of the fine-tuned BERT model the questions generated beforehand using both the source article and the summary to get two sets of answers.
 - 6: We compare the corresponding answers using an answer similarity metric, and we get a final score averaging the answer similarity metric (f1 score) over all questions
-

- QAGS




Figure 4.4: Schema of QAGS

4.2.2 SUMMAQA

SUMMAQA works similarly to QAGS, the main difference is that there is not a generative model to create the questions, in particular:

SUMMAQA

- 1: We find all the entities in the source text (the original ticket) with Spacy
 - 2: Create one sentence (which will be our question) for each entity. We keep only the sentence in which the entity is if there are more sentences. The entity will be masked. Ex: "Yesterday I went to Paris" becomes "Yesterday I went to [MASK]"
 - 3: The MASKED entities will be the True labels when calculating the metrics
 - 4: We answer the questions generated before using a fine-tuned version of BERT model, using both the source and the summarization/anonymization as contexts
 - 5: We compare the answers and calculate the f1 score for both contexts
-

4.2.3 Cosine Similarity

We measured the cosine similarity between sentence embeddings of original text and summarization. The sentence embeddings used were calculated using Sentence-BERT, a modified version of the pretrained BERT network that uses siamese networks to obtain more meaningful embedding of the sentence, compared to getting the embedding of the [CLS] token or to average all the other embeddings.

Result

After testing all of the three methods described before, the one chosen was QAGS. The main reason why is its flexibility. Right now it was not implemented, but in future works it could be possible to add to the questions automatically generated by BART some questions written by the users to check whether a piece of information has been anonymized or not.

However, the numeric results are difficult to interpret, because sometimes we want to retain some information from the original text and sometimes we would prefer the information to be removed.

- SUMMAQA




Figure 4.5: Schema of SUMMAQA

COSINE SIMILARITY

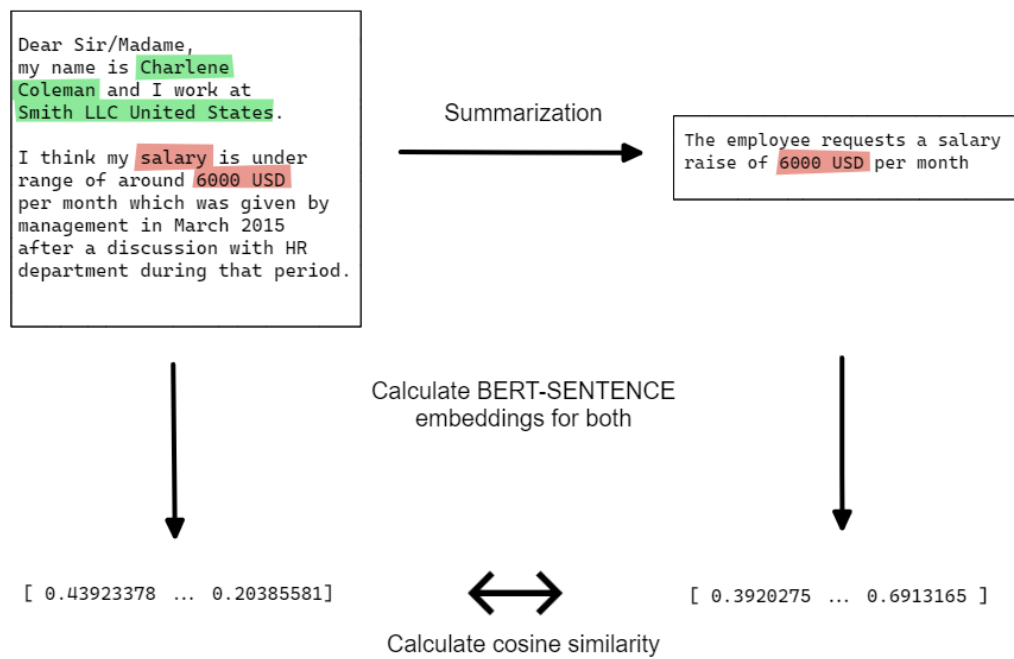


Figure 4.6: Schema of Cosine Similarity for

4.3 Named Entity Recognition

Named entity recognition is the process of automatically identifying and classifying named entities in a text. This can include identifying and categorizing named entities such as people, organizations, locations, and so on.

NER systems might be used to automatically extract information from a large collection of documents, such as identifying all mentions of specific named entities or analyzing the relationships between different named entities.

4.3.1 Weak labeling

During the creation of the ticket, we tried to automatically save also the entities of the ticket. With entities, we do not mean the classical entities used for NER, but the original variables specific to a ticket's category (complete list in Table 3.1).

Finding the entities in the prompt is trivial, since the positions are fixed in the template, and the filling operation is managed by us.

Example:

```
From: ${email}
To: ${company email}
First name: ${first name}
Last name: ${last name}
Company: ${company}
Date: ${ticket date} Date start absence: ${date_start_absence}
Reason absence: ${reason_absence}
Subject: Request for sick leave for ${days}
```

Though the template and prompt serve as our starting point to generate the actual ticket text, it is only this latter part that would be present in a real ticket. Identifying the locations of the entities is not straightforward since I usually provide only the initial information and a limited template, with the rest of the generation being done by GPT.

Once the ticket is generated, to find the entities in the generated text I use 2 different approaches:

- Exact match
- Heuristics

The exact match is straight-forward, it just looks for the exact same words in the generated text, and if the word/words are found, the entity is added to the list of entities of the ticket. However this method often does not work, since GPT could use only part of the entity (Ex: "A medical consultation" is in the text as "a consultation") or slightly alter the entity (ex: 01/10/1998 → the first of October of 1998)

This is the reason why we implemented manually a set of different rules to find modified versions of the original entity in the text. These versions do not cover every possible case, since they are hard-coded by us from experience and by looking at how GPT behaves, this is why we called this approach "Heuristics".

Some of these heuristics are:

- Check not only the full entity, but also subsets of it removing stopwords
Example:
reason_of_absence: "An Urinary Tract Infection" →
 - "Urinary Tract Infection"
 - "Urinary Tract"
 - "Tract Infection"
 - "Urinary"
 - "Tract"
 - "Infection"
- Check all the possible version of a percentage text
"5%", "5 %", "5.0%", "5 percent", ...
- Check different formats of date
MM/DD/YYYY, DD/MM/YYYY, "First of October", "1st of October", ...

With this method, we were able to identify entities in only 6272 tickets out of the total 16000 tickets. The complete list of all the entities found is shown in Table 4.3.

duration	1016	increase_in_percentage	539
location	1411	work_title	33
date	70	wage_gap	421
work_shift	70	number_of_days	982
reason_of_change	7	date_start_absence	3
to_who	1221	description_life_event	1214
reason	1119	date_travel	112
complaint	106	airport	28
salary	376		

Table 4.3: Entities found with heuristics

4.3.2 Classical NER

Token classification is a common technique used in NER to identify and classify named entities in a text. In token classification, the text is first divided into individual tokens, which are typically words or phrases. The model then predicts a label for each token, indicating the type of named entity it represents.

To assign the labels to the tokens we used the BIO(Beginning, Inside, Outside) format, where the B tag is used for the first token of an entity, the I tag for the rest of the tokens of the entity and lastly the O tag is used for the tokens that do not belong to any entities.

Example:

Alex B-PER
is O
going O
to O
Los B-LOC
Angeles I-LOC
in O
California B-LOC

The labels assigned to each token are then used for a token classification task, using a RoBERTa model and the Spacy training pipeline.

We used the default Spacy parameters for the training.

The survey tickets, used as always as our test set, were manually labeled by ourselves.

The results were not satisfying, we achieved a f_1 score of 0.34 on the exact matching and a f_1 score of 0.45 on the partial matching. When the model finds the correct entity for all the tokens and only the tokens of an entity, it's considered an exact matching. Whereas, when the model finds the correct entity for only a subset of the token of an entity, then is considered a partial match

Example exact match:

to O
 Los LOC
 Angeles LOC
 in O

Example partial match:

to O
 Los LOC
 Angeles O
 in O

The main problem we encountered was that the NER model was not able to recognize the different entities of the same type depending on the context. For example, the model could not recognize the difference between a date of start absence (entity of request of time off due to health reasons) and a date of travel (entity of request of refund for travel).

4.3.3 Our approach to NER

To overcome the limits of the first approach, we developed a new approach based on transfer learning and sentence classification.

The main intuition behind the new approach is to use the pre-trained spacy NER model and exploit the entities found by the base model to extract our entities.

First of all, we build a dictionary of the matchings between entities of Spacy and our entities. We decided to work on a subset of our entities and not consider all the entities that did not have a clear matching with Spacy entities. In Table 4.4 we show all the matchings. Then, we scan our dataset and look

Category	Entities matching
Ask information_Accommodation	{"location": "GPE", "duration": "DATE"}
Life event_Health issues	{"date_start_absence": "DATE", "number_of_days": "DATE"}
Refund_Refund travel	{"date_travel": "DATE", "location": "GPE"}
Salary_Gender pay gap	{"wage_gap": "PERCENT"}
Salary_Salary raise	{"increase_in_percentage": "PERCENT", "salary": "MONEY"}
Timetable change_Shift change	{"date": "DATE"}

Table 4.4: Matching of our entities with Spacy entities

for all categories, the entities shown in Table 4.4. If we find an entity not present in the table, it is filtered out. For each entity, we will have an array of the type [ENTITY, START_CHAR, END_CHAR] (Figure 4.7).

After that, we tokenize the text of all the tickets and we look up for each ticket its entities' tokens. Then for each entity, we save the indexes of the corresponding tokens, so for example if the entity "twelve months" is split into two tokens "twelve" and "month", which are the 22nd and 23rd tokens of the text, then we will save ["twelve months", DATE, [22,23]].

We also had to manage the case in which a token is not present in the tokenizer vocabulary, so is split into sub-tokens. We took advantage of the fact that if a word is split into two or more words the new tokens will have the prefix "##".

The number of tokens for each entity can vary, for practical reasons in the training phase we set the maximum number of tokens to $N = 50$ (Figure 4.8). Now we have all we need for the training. The base architecture is a *bert-base-cased* model, which we modified on the last layer.




Figure 4.7: Ticket NER preprocessing part 1




Figure 4.8: Ticket NER preprocessing part 2

In the training phase, we pass to the BERT model at each step the tickets' texts and their entities with the indexes of the matching tokens. For each entity of the ticket, there will be a training record. The ticket will be unpacked into 512 tokens, and each token will be processed by the BERT architecture (same architecture as the classifier Figure 4.1).

At the last layer, instead of passing the [CLS] token to a classifier, we concatenate the [CLS] token with the average of all the embeddings of the tokens of the current entity. This is the reason why we preprocessed the dataset to find the indexes of the tokens' indexes for the entities. In BERT base each token has a dimension $d = 768$, so the input of the last feed forward

neural network will be $N = 1536$. The output dimension will be $M = 10$ (Figure 4.9). The labels of the training are the combination of the ticket category and the entity, since for each entity found in a ticket we generate a record (For example if there are 3 entities in a ticket, there will be 3 records in the training set with the same ticket text, but that will have a different token embedding for the classification in the last layer and they will have different labels)

At test time, for each ticket in the test set, we analyze it with Spacy and




Figure 4.9: Ticket NER training

we find all the entities. Then we filter the entities: we cannot filter based on the label, since at test time in theory we do not know the label of the tickets. Therefore we filter the entities based on the list of all entities used in all tickets' categories.

Then, as in the training phase, for each entity we build a new record composed of the output of the [CLS] token through the model architecture and of the average of the tokens' embeddings.

For each ticket in the test set we obtain a prediction of its category combined with the entity currently analyzed.

Results

We are able to achieve way better results with the new method that combines a ticket classifier with a NER model. On the test set we achieve a f_1 score of 0.96.

In Figure 4.10 we show the confusion matrix of the classification.

We believe that we are able to achieve such good results because there are few entities that are shared between more categories, and the union of the information of the ticket ([CLS] token) and the information of the entity makes it easy for the model to distinguish both the category and the entity.




Figure 4.10: Ticket NER confusion matrix

Chapter 5

Ethical Considerations

Building a dataset to train machine learning models, it is important to consider the possible ethical consequences because the data that a model is trained on can influence the model’s behavior and decisions. If the data is biased or unrepresentative, the model’s predictions and decisions may also be biased or unrepresentative. This can lead to unfair or unjust outcomes for individuals or groups of people.

We are aware that our personal biases could have influenced the topics and the realities described by the tickets that compose the dataset. We have tried to make the dataset as balanced as possible in terms of gender.

In some cases, we inferred prompts to create situations that talk about discrimination at workplaces based on gender, race, disability, religion, or nationality. In these cases, there are several tickets that contain toxic language. Nevertheless, Language Models can create unfair discrimination by perpetuating stereotypes and social biases [31][32]. Most of our work is based on the use of one of these LMs, so we cannot guarantee that there are no biases due to the original data on which the generative model was trained. Some of these biases result in underrepresented groups, stereotypes, exclusionary norms, unfair discrimination and lower performance by social group[33].

Chapter 6

Conclusion

Firstly, we developed a taxonomy of HR tickets and created a series of templates for each category. We discussed the considerations behind constructing the templates and then researched free-open datasets related to each category. We employed privacy-preserving techniques such as Bayesian Networks to process the datasets.

Then, we identified the state-of-the-art models for text generation, followed by a discussion of their respective novelties. After that, we discussed the generative model chosen by us, GPT-J, and the reasonings behind the decision. We explained in detail the theory behind Transformers and the novelties of GPT-J with respect to the other generative models.

We examined the model's ability to generate texts from two angles: we looked how the inputs influence the outputs and how each neuron of the model behaves across different steps of the generation.

We conducted a survey in order to demonstrate that our dataset could resemble real tickets and that the tickets created by our application could be used to train machine learning models which worked also in real scenarios. Specifically we tested the usefulness of the dataset on three different use cases: classification of tickets, anonymization of tickets and named entity recognition.

For the classification of tickets we trained and tested two different models. Whereas for the anonymization task, we used a novel approach that summarizes and rewrites the original text. Lastly, we experimented a classical named entity recognition model on our dataset and, after achieving not satisfying

results, we developed a new technique that unifies a text classification task with named entity recognition.

We achieved good results on all three different tasks, therefore we can confidently state that the dataset we built could be used to train ML models for real-world applications. We recognize that each company has a distinct taxonomy for tickets, so we developed the *Ticket Generator* application to facilitate the adaptation to different contexts.

Bibliography

- [1] SAP. *SAP Security Research*. <https://www.sap.com/documents/2022/03/8c5580cc-207e-0010-bca6-c68f7e60039b.html>. [Online; accessed]. 2022.
- [2] European Commission. *Regulation (EU) 2016/679 of the European Parliament and of the Council of 27 April 2016 on the protection of natural persons with regard to the processing of personal data and on the free movement of such data, and repealing Directive 95/46/EC (General Data Protection Regulation) (Text with EEA relevance)*. 2016. URL: <https://eur-lex.europa.eu/eli/reg/2016/679/oj>.
- [3] Steven M Bellovin, Preetam K Dutta, and Nathan Reiter. «Privacy and synthetic datasets». In: *Stan. Tech. L. Rev.* 22 (2019), p. 1.
- [4] Ildikó Pilán, Pierre Lison, Lilja Øvrelid, Anthi Papadopoulou, David Sánchez, and Montserrat Batet. «The text anonymization benchmark (tab): A dedicated corpus and evaluation framework for text anonymization». In: *arXiv preprint arXiv:2202.00443* (2022).
- [5] Yi-Chun Chen, Tim A Wheeler, and Mykel J Kochenderfer. «Learning discrete Bayesian networks from continuous data». In: *Journal of Artificial Intelligence Research* 59 (2017), pp. 103–132.
- [6] Jun Zhang, Graham Cormode, Cecilia M Procopiuc, Divesh Srivastava, and Xiaokui Xiao. «Privbayes: Private data release via bayesian networks». In: *ACM Transactions on Database Systems (TODS)* 42.4 (2017), pp. 1–41.
- [7] Daniele Faraglia and Other Contributors. *Faker*. URL: <https://github.com/joke2k/faker>.

- [8] Alec Radford, Karthik Narasimhan, Tim Salimans, Ilya Sutskever, et al. «Improving language understanding by generative pre-training». In: (2018).
- [9] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. «Language models are unsupervised multitask learners». In: *OpenAI blog* 1.8 (2019), p. 9.
- [10] Tom Brown et al. «Language models are few-shot learners». In: *Advances in neural information processing systems* 33 (2020), pp. 1877–1901.
- [11] Tom Brown et al. «Language models are few-shot learners». In: *Advances in neural information processing systems* 33 (2020), pp. 1877–1901.
- [12] Aakanksha Chowdhery et al. «Palm: Scaling language modeling with pathways». In: *arXiv preprint arXiv:2204.02311* (2022).
- [13] Jordan Hoffmann et al. «Training Compute-Optimal Large Language Models». In: *arXiv preprint arXiv:2203.15556* (2022).
- [14] Ben Wang and Aran Komatsuzaki. *GPT-J-6B: A 6 Billion Parameter Autoregressive Language Model*. <https://github.com/kingoflolz/mesh-transformer-jax>. May 2021.
- [15] Teven Le Scao et al. «BLOOM: A 176B-Parameter Open-Access Multilingual Language Model». In: *arXiv preprint arXiv:2211.05100* (2022).
- [16] Leo Gao et al. «The pile: An 800gb dataset of diverse text for language modeling». In: *arXiv preprint arXiv:2101.00027* (2020).
- [17] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. «Layer normalization». In: *arXiv preprint arXiv:1607.06450* (2016).
- [18] Jingjing Xu, Xu Sun, Zhiyuan Zhang, Guangxiang Zhao, and Junyang Lin. «Understanding and improving layer normalization». In: *Advances in Neural Information Processing Systems* 32 (2019).
- [19] LucyTalks. *The Modern Tokenization Stack for NLP: Byte Pair Encoding*. <https://lucytalksdata.com/the-modern-tokenization-stack-for-nlp-byte-pair-encoding>. [Online; accessed]. 2022.
- [20] Jianlin Su, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. «Roformer: Enhanced transformer with rotary position embedding». In: *arXiv preprint arXiv:2104.09864* (2021).

- [21] Stella Biderman, Sid Black, Charles Foster, Leo Gao, Eric Hallahan, Horace He, Ben Wang, and Phil Wang. *Rotary Embeddings: A Relative Revolution*. blog.eleuther.ai/. [Online; accessed]. 2021.
- [22] J Alammam. «Ecco: An Open Source Library for the Explainability of Transformer Language Models». In: *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing: System Demonstrations*. Association for Computational Linguistics, 2021.
- [23] Avanti Shrikumar, Peyton Greenside, and Anshul Kundaje. «Learning important features through propagating activation differences». In: *International conference on machine learning*. PMLR. 2017, pp. 3145–3153.
- [24] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. «Axiomatic attribution for deep networks». In: *International conference on machine learning*. PMLR. 2017, pp. 3319–3328.
- [25] Emily Pitler and Ani Nenkova. «Revisiting readability: A unified framework for predicting text quality». In: *Proceedings of the 2008 conference on empirical methods in natural language processing*. 2008, pp. 186–195.
- [26] Scott A Crossley, Jennifer L Weston, Susan T McLain Sullivan, and Danielle S McNamara. «The development of writing proficiency as a function of grade level: A linguistic analysis». In: *Written Communication* 28.3 (2011), pp. 282–311.
- [27] Danielle S McNamara, Scott A Crossley, and Philip M McCarthy. «Linguistic features of writing quality». In: *Written communication* 27.1 (2010), pp. 57–86.
- [28] Maarten Grootendorst. «BERTopic: Neural topic modeling with a class-based TF-IDF procedure». In: *arXiv preprint arXiv:2203.05794* (2022).
- [29] Rohola Zandie and Mohammad H Mahoor. «Topical language generation using transformers». In: *Natural Language Engineering* (2021), pp. 1–23.
- [30] Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, and Luke Zettlemoyer. «Bart: Denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension». In: *arXiv preprint arXiv:1910.13461* (2019).

- [31] Abubakar Abid, Maheen Farooqi, and James Zou. «Persistent anti-muslim bias in large language models». In: *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*. 2021, pp. 298–306.
- [32] Li Lucy and David Bamman. «Gender and Representation Bias in GPT-3 Generated Stories». In: *Proceedings of the Third Workshop on Narrative Understanding*. Virtual: Association for Computational Linguistics, June 2021. URL: <https://aclanthology.org/2021.nuse-1.5>.
- [33] Laura Weidinger et al. «Ethical and social risks of harm from language models». In: *arXiv preprint arXiv:2112.04359* (2021).